



Termal Defect Detection

Systems and Architectures for Big Data

A.Y. 2024/2025

Martorelli Luca
0357263

Iurato Chiara
0341741

Cortese Alessandro
0350035

Agenda

01

INTRODUCTION

02

SYSTEM ARCHITECTURE

03

PIPELINE

04

PROCESSING QUERIES

05

KAFKA SYSTEM ARCHITECTURE

06

KAFKA STREAMS PIPELINE

07



RESULTS

08

CONCLUSIONS

Introduction

- L-PBF (Laser Powder Bed Fusion) is a metal 3D printing technique that fuses powder layer by layer using a laser, enabling design flexibility and minimizing material waste.
- However, it can produce critical defects (e.g., porosity) often detected only after production, leading to wasted time and resources.
- To address this, Optical Tomography (OT) enables real-time monitoring by capturing thermal images layer by layer during printing.
- OT images reveal temperature distributions, allowing early detection of thermal anomalies and timely interventions.
- Each data batch is a stream of high-resolution 16-bit TIFF OT images, divided into tiles representing local blocks of each printed layer.



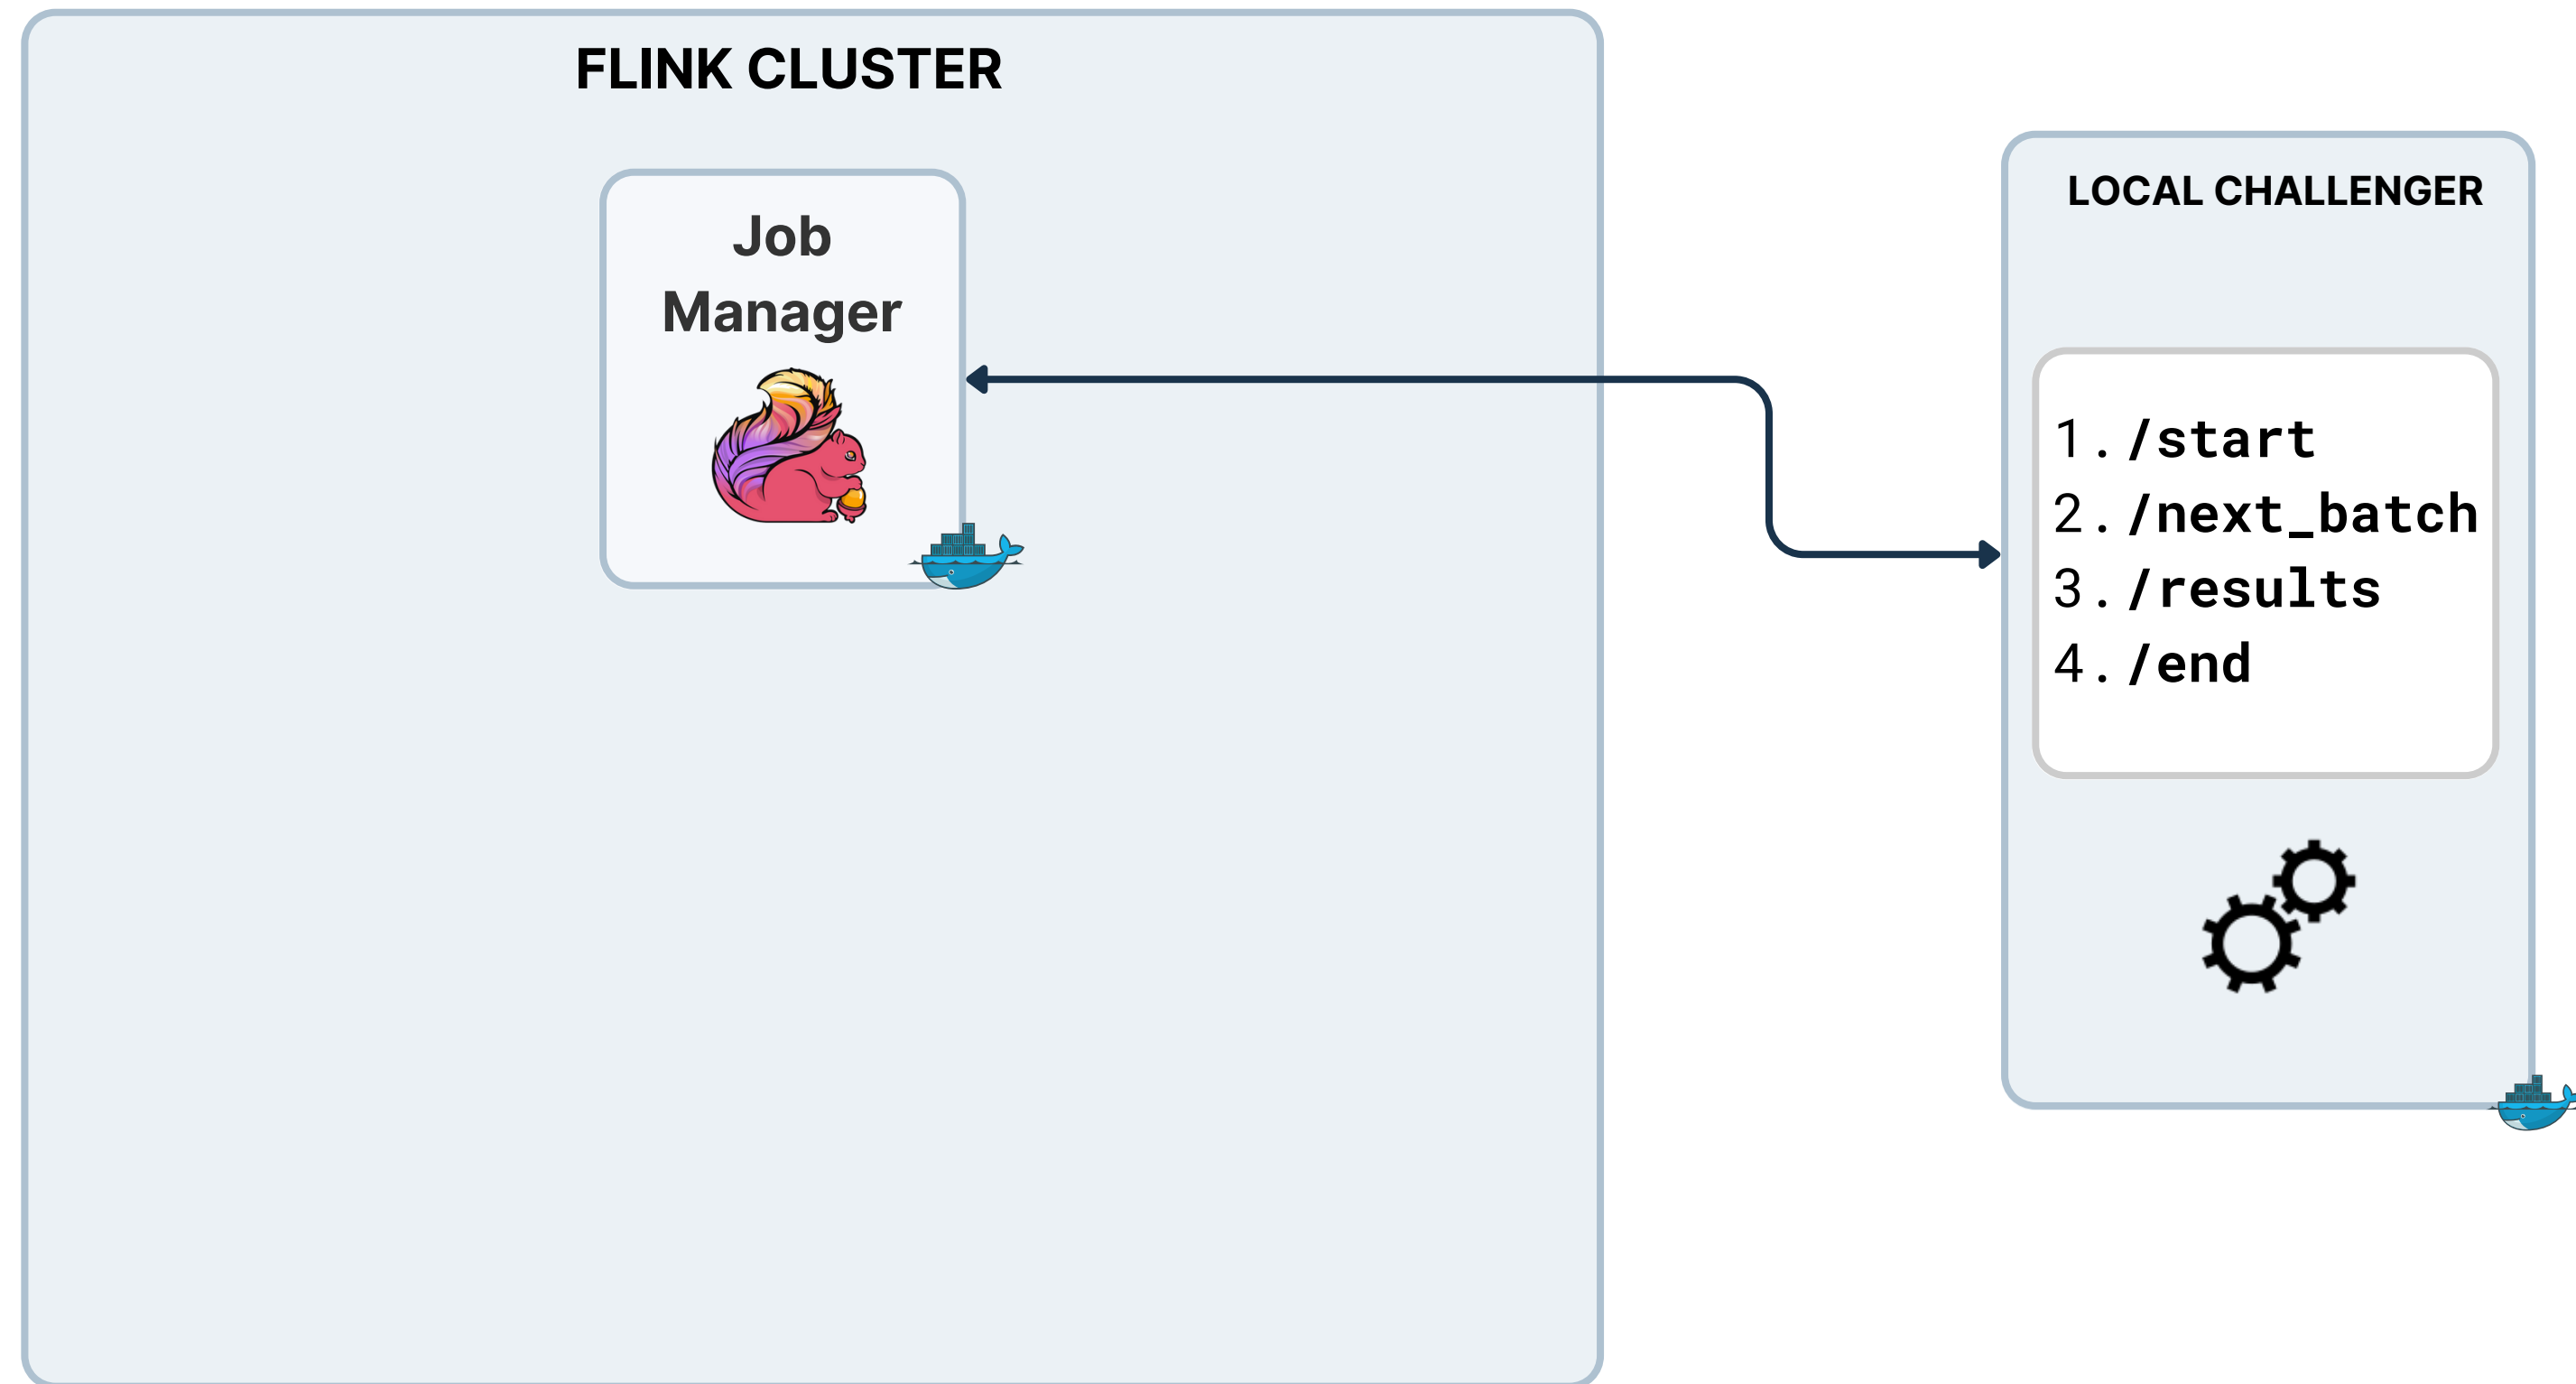
Flink System Architecture

LOCAL CHALLENGER

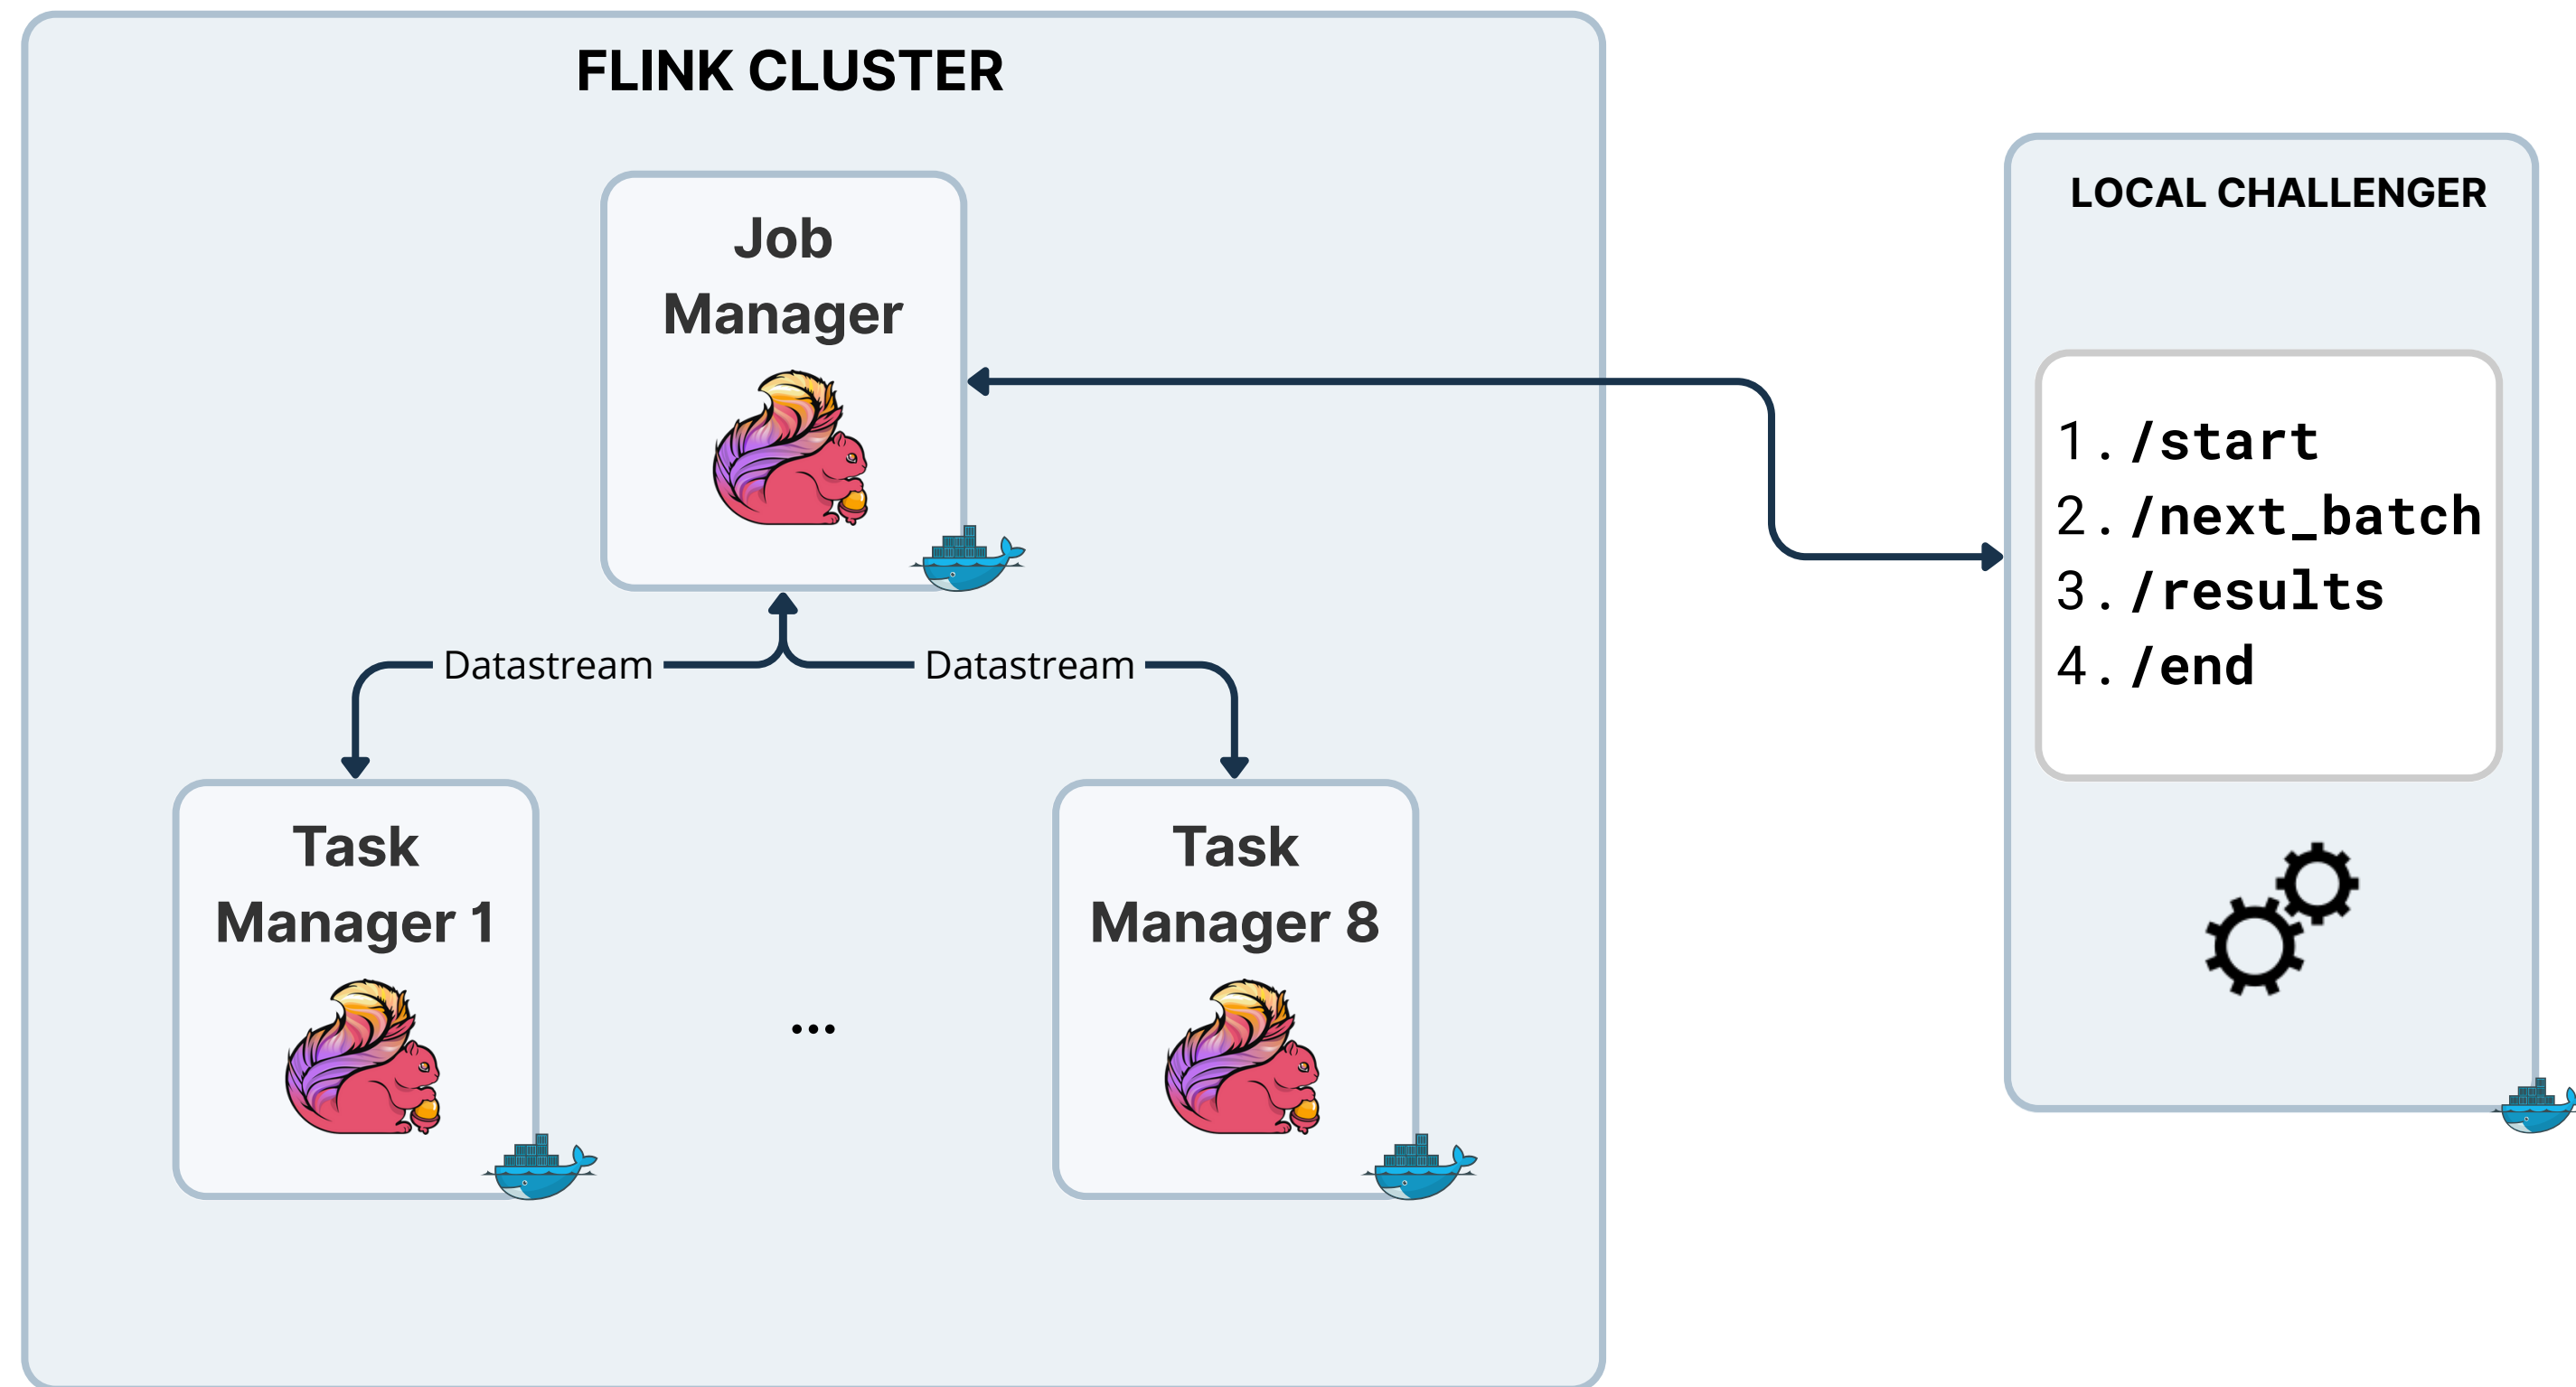
1. `/start`
2. `/next_batch`
3. `/results`
4. `/end`



Flink System Architecture

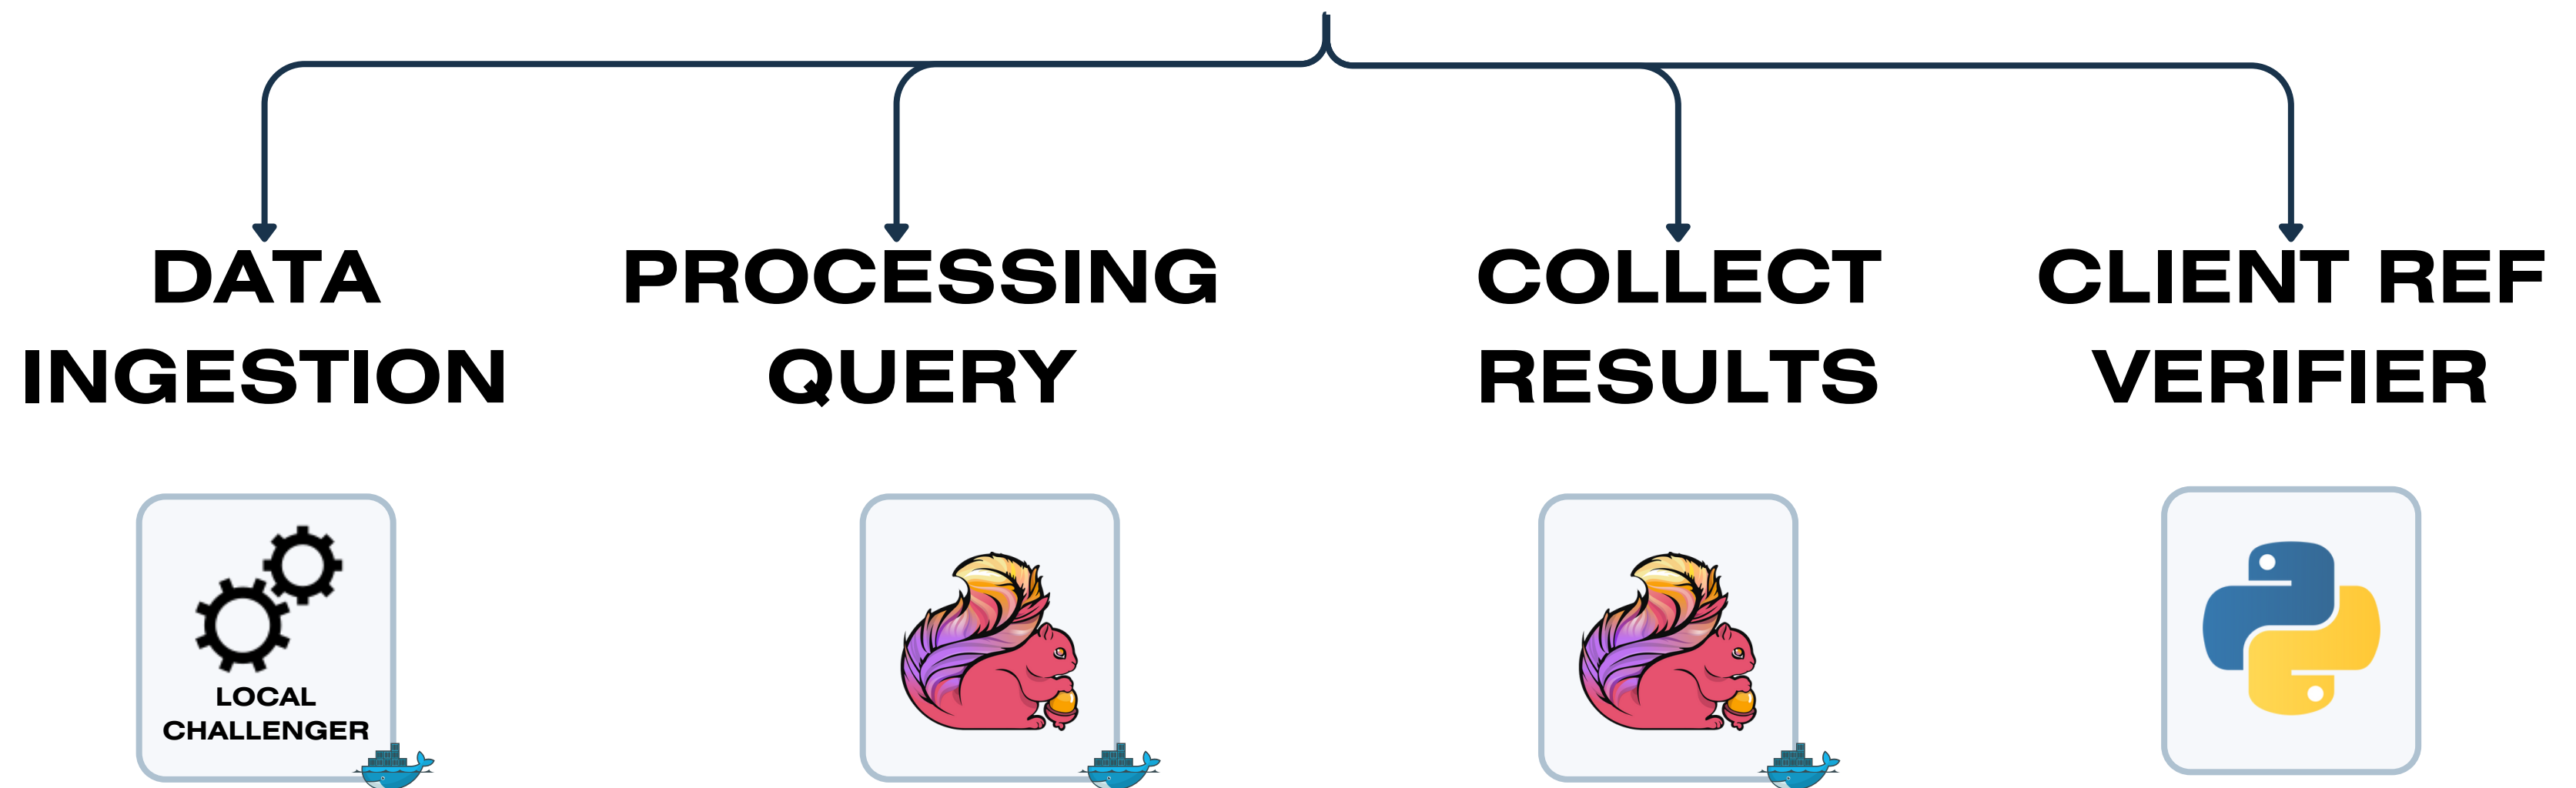


Flink System Architecture

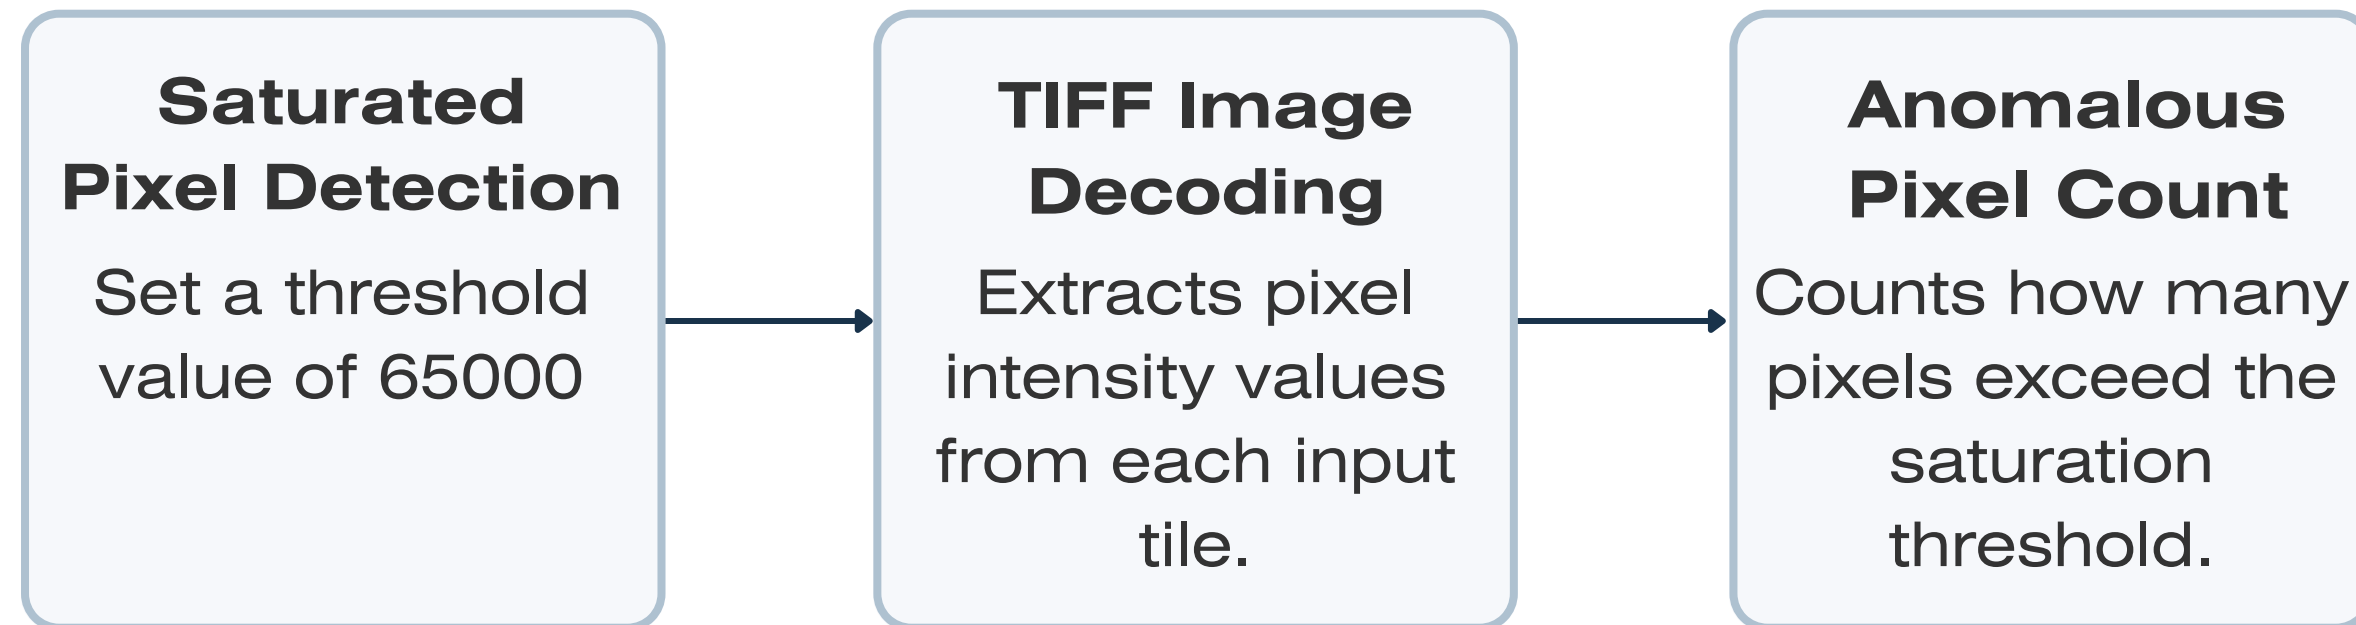


Flink Pipeline

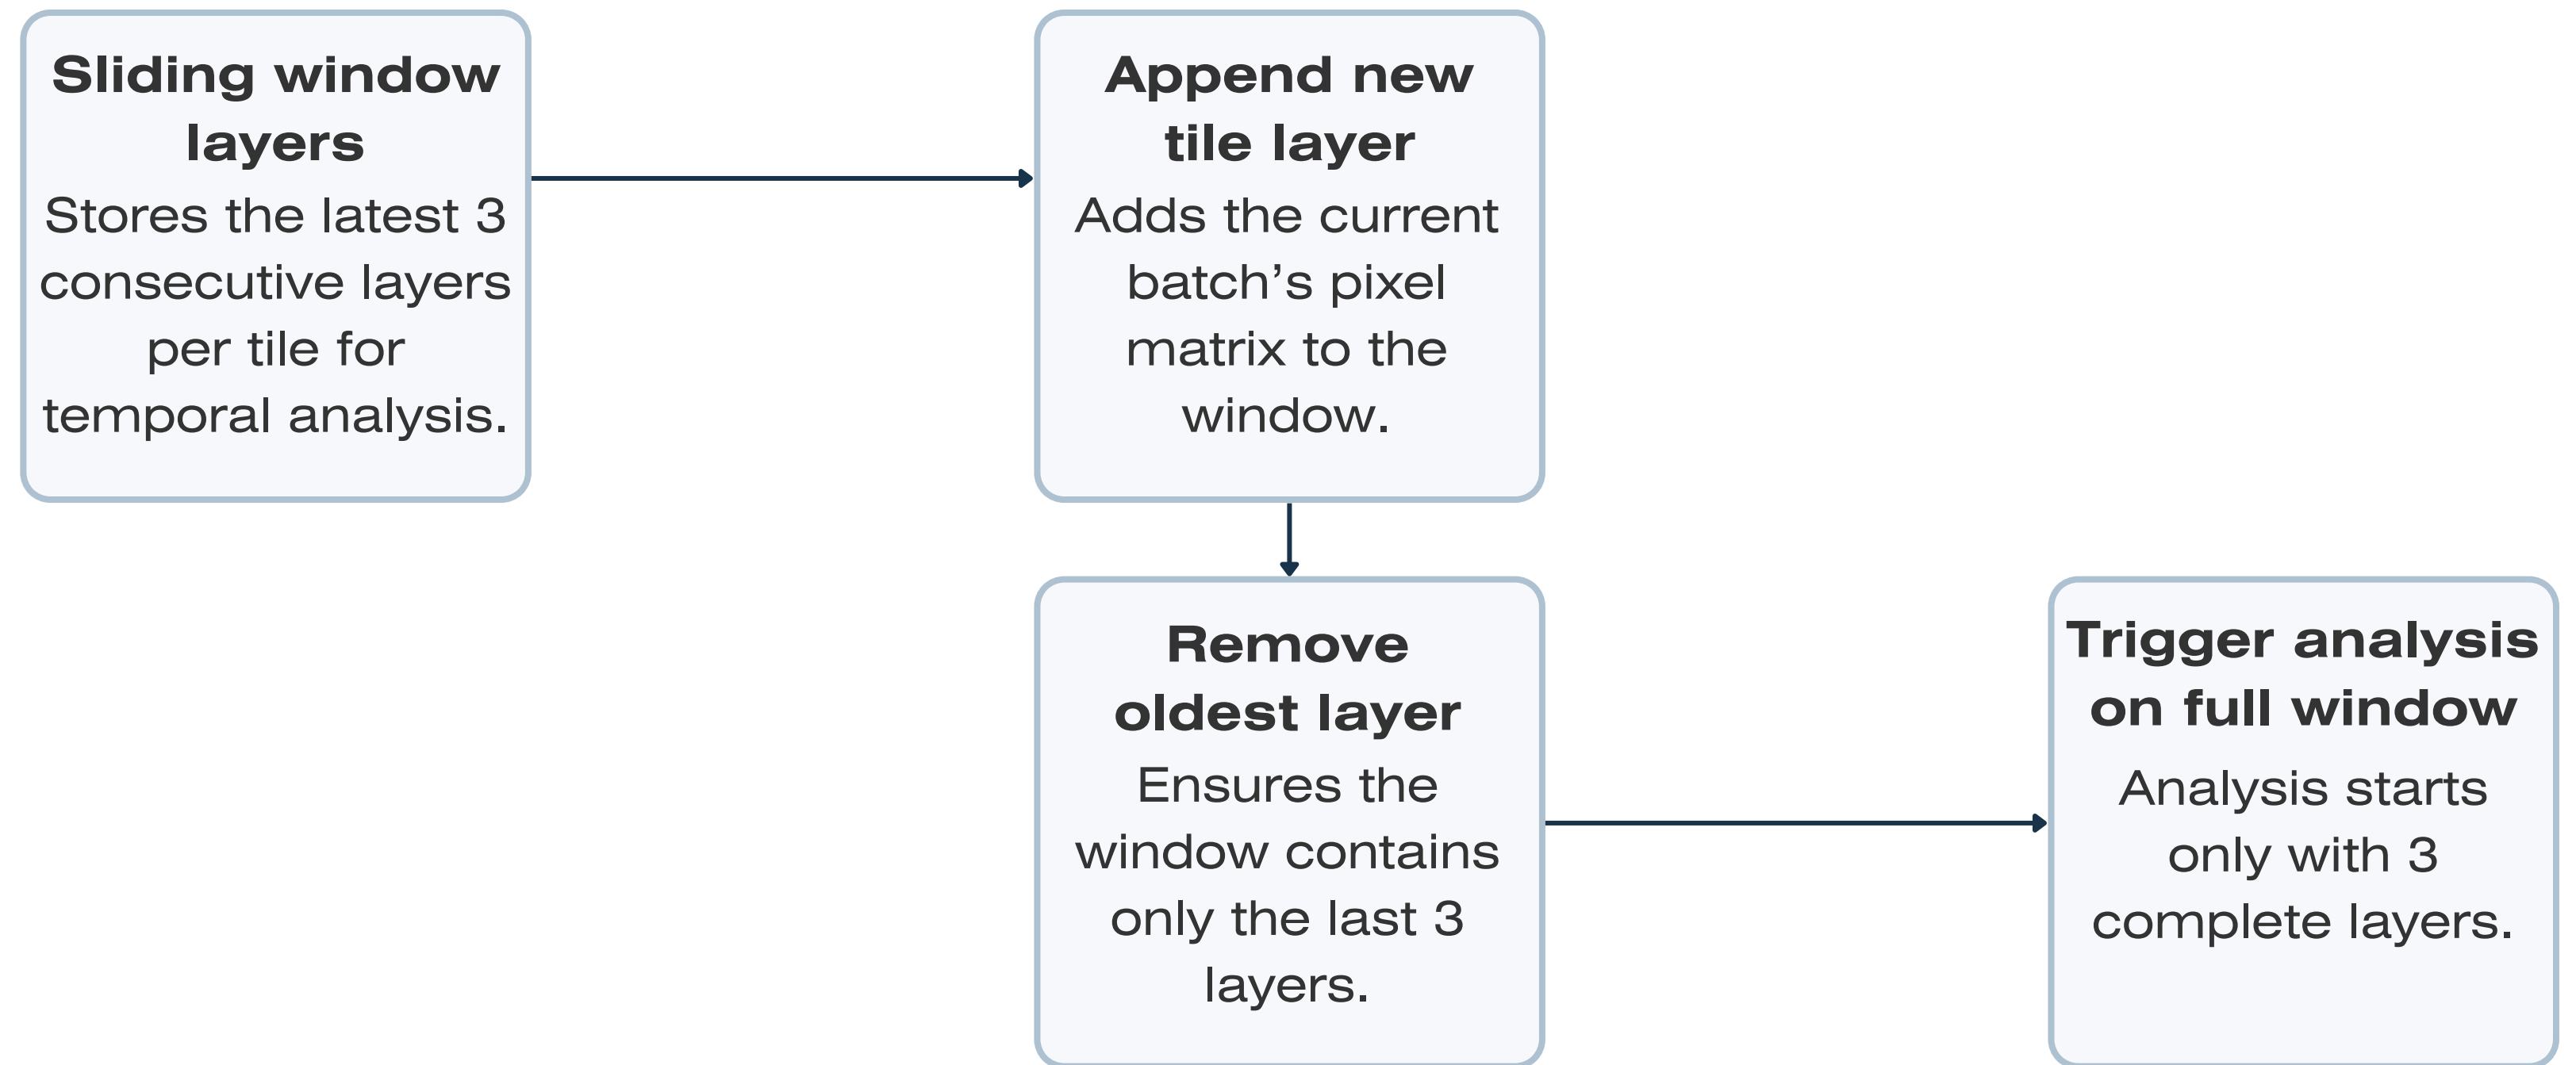
The data processing pipeline consists of the following steps:



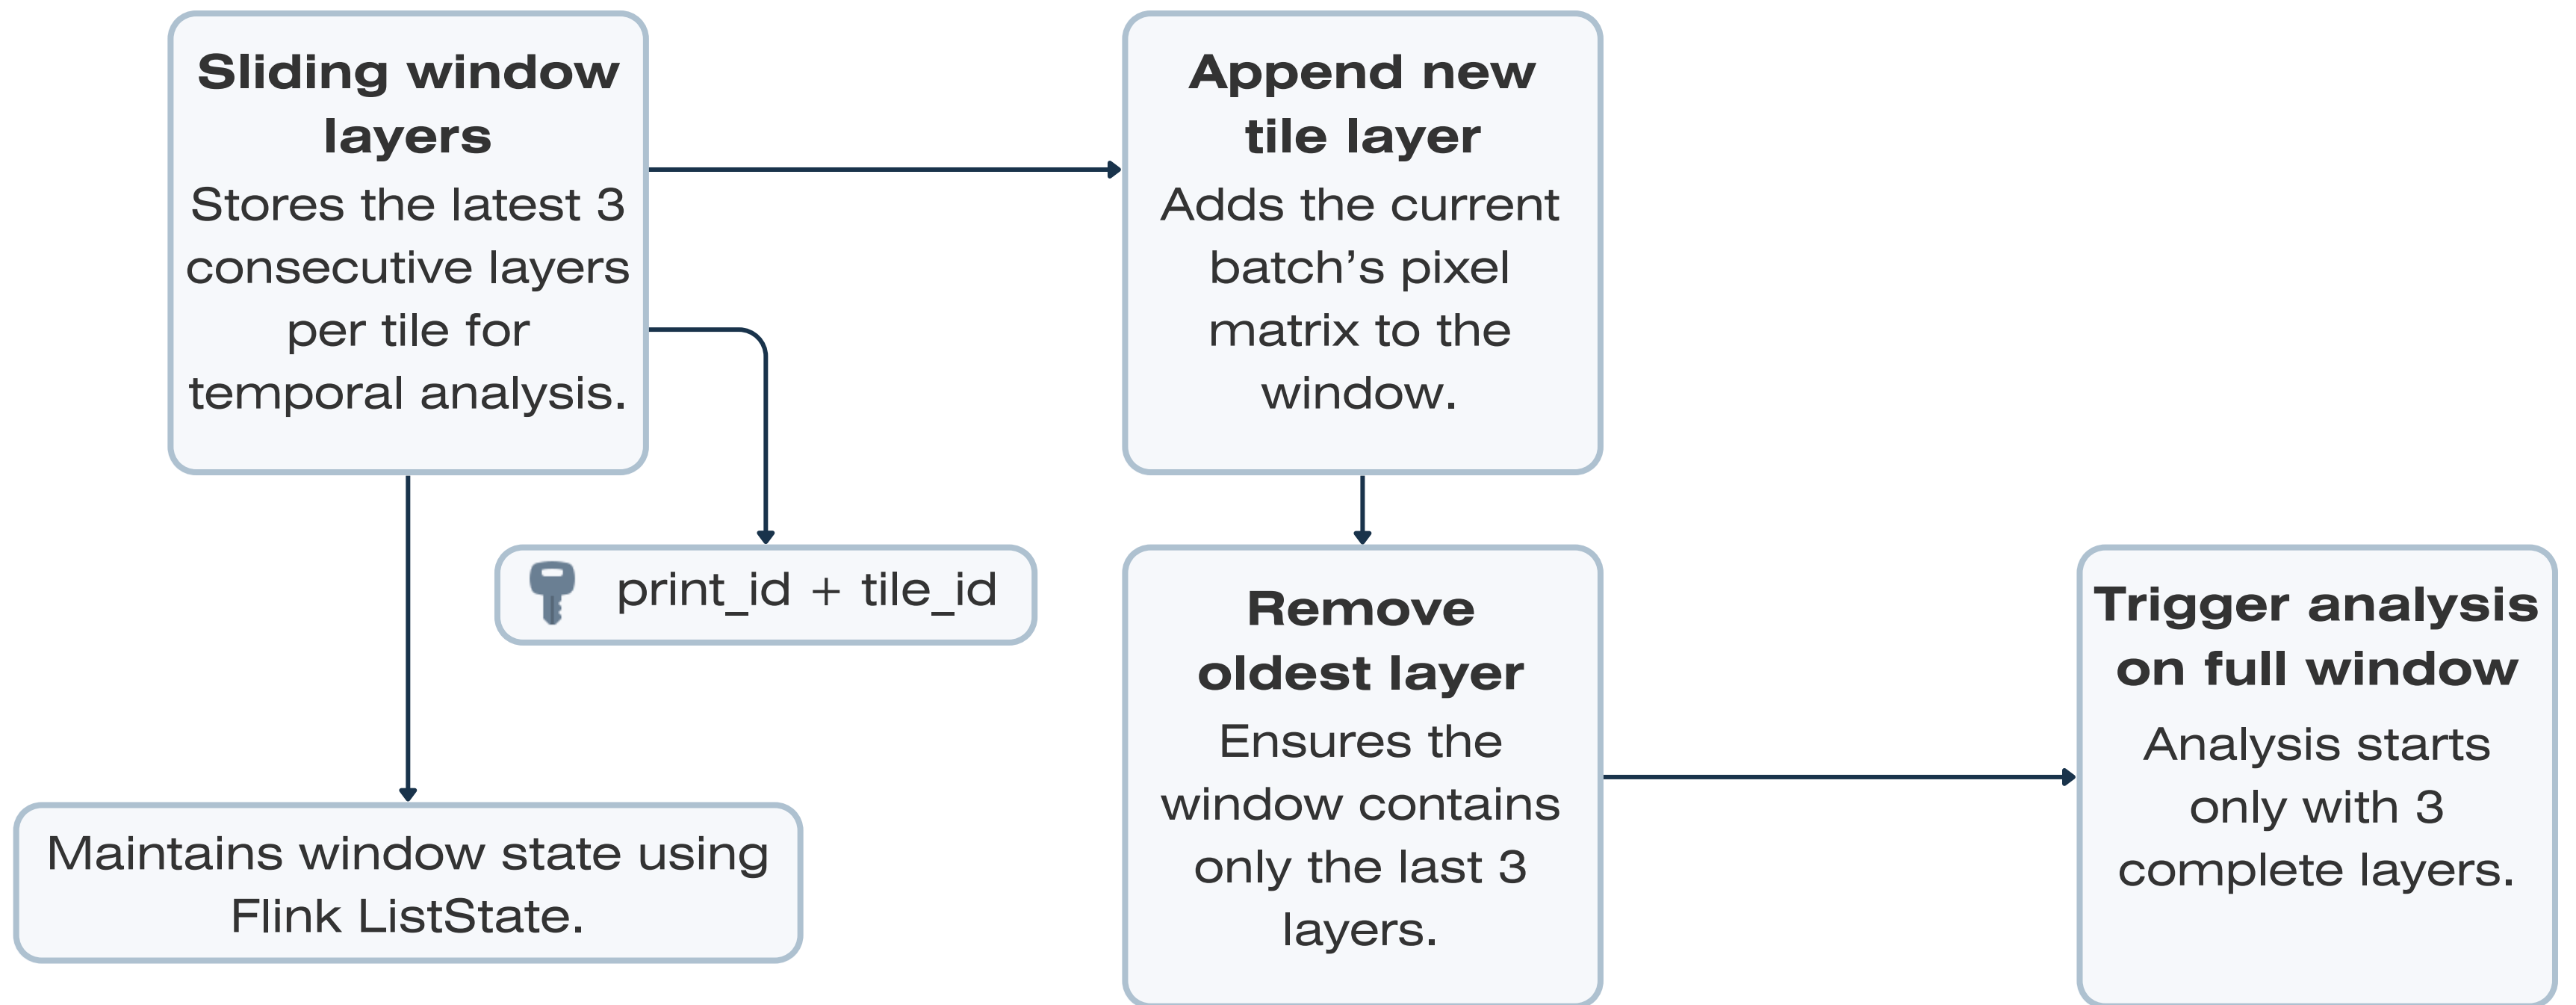
Query 1 - Flink



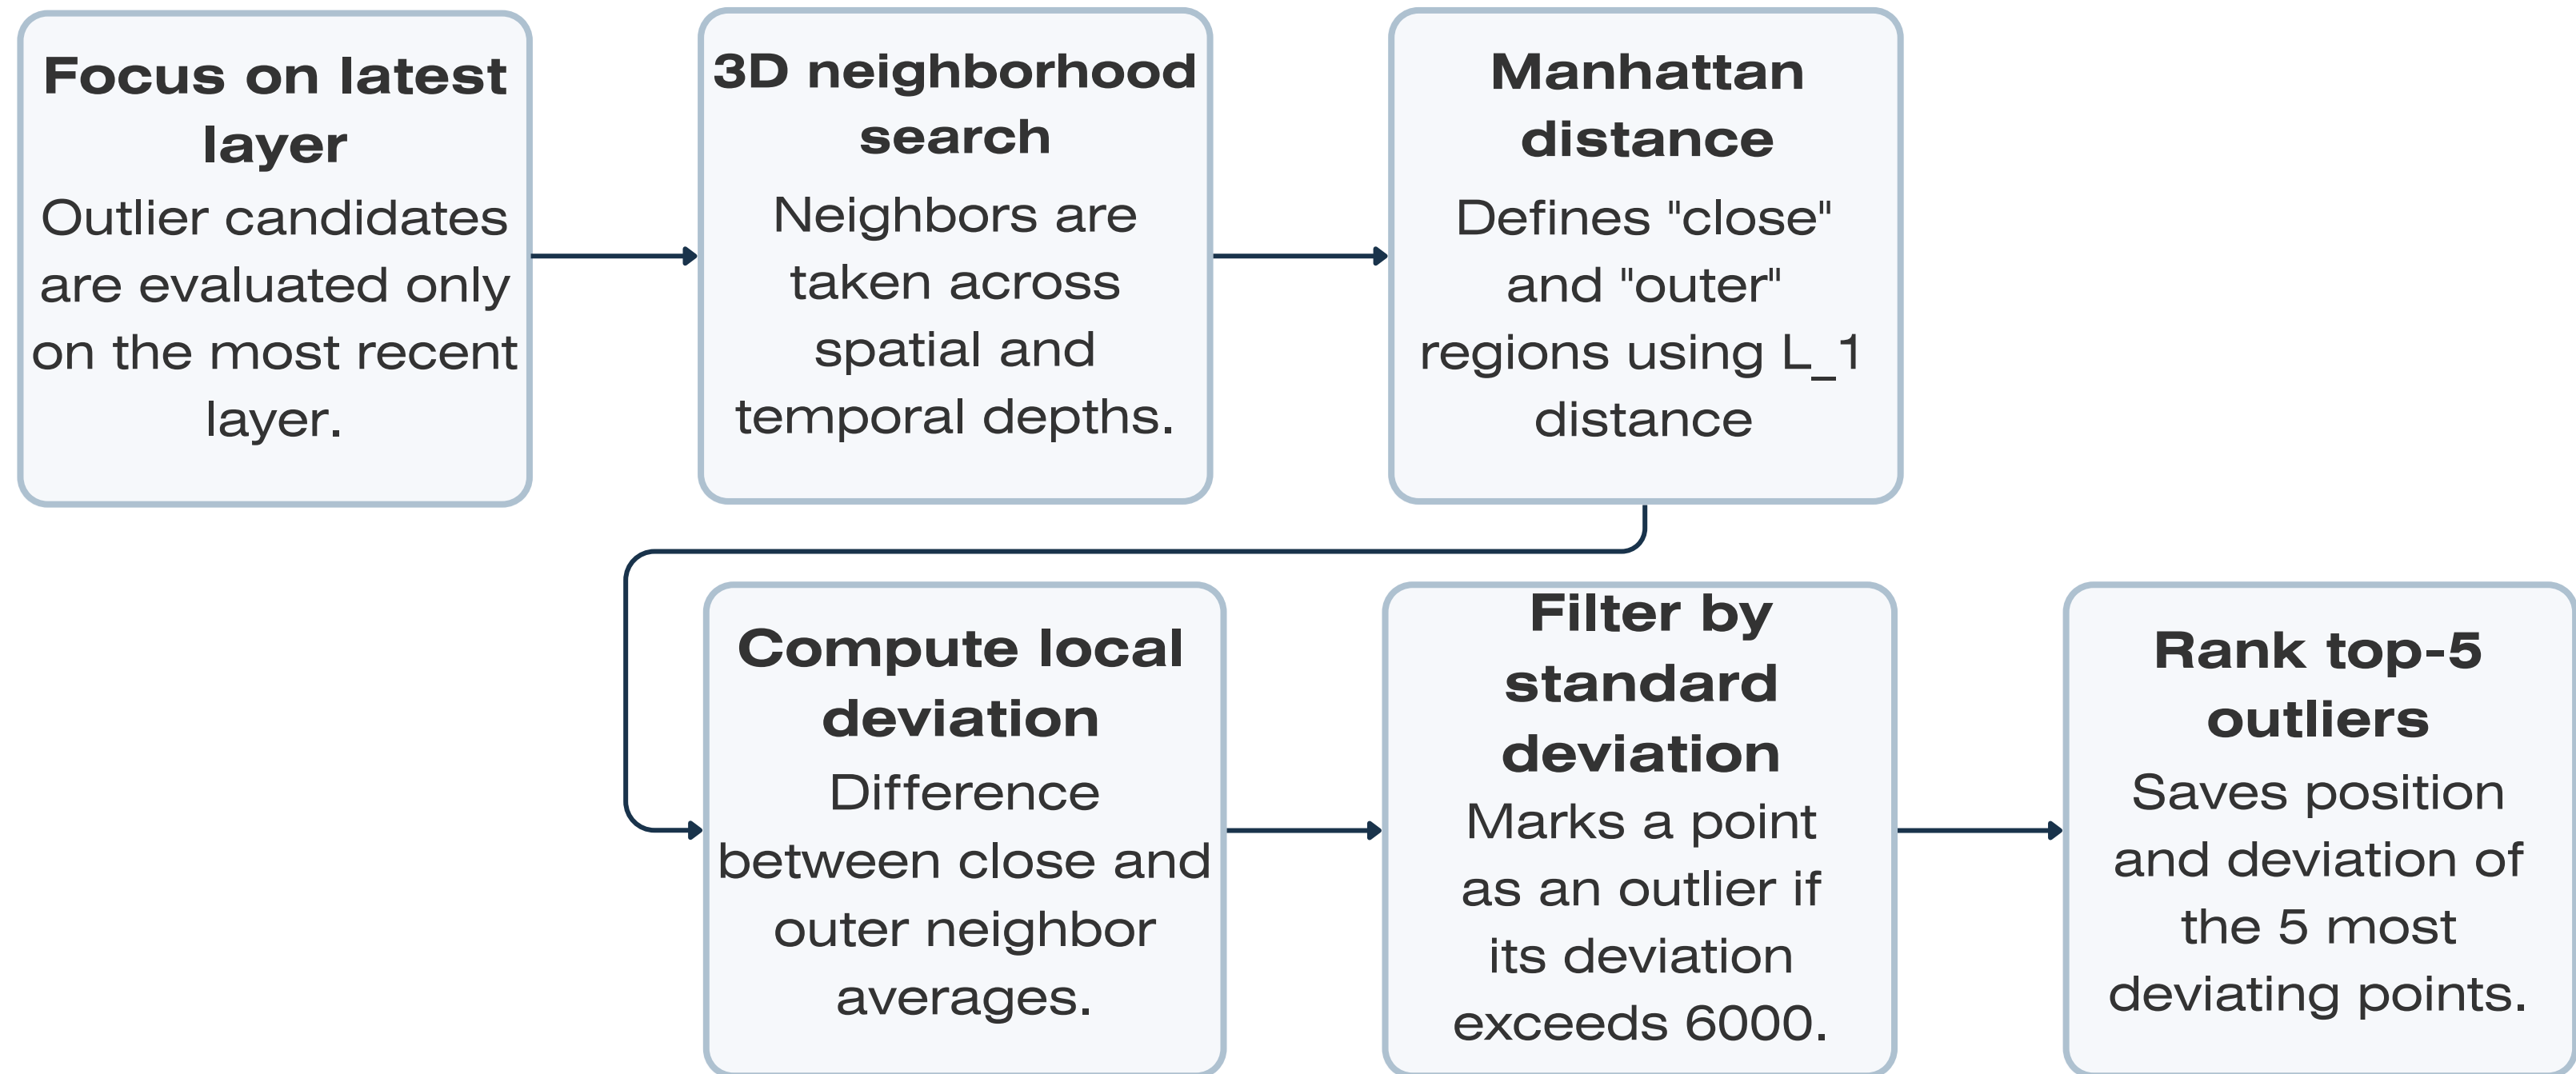
Query 2 - Flink (Windowing)



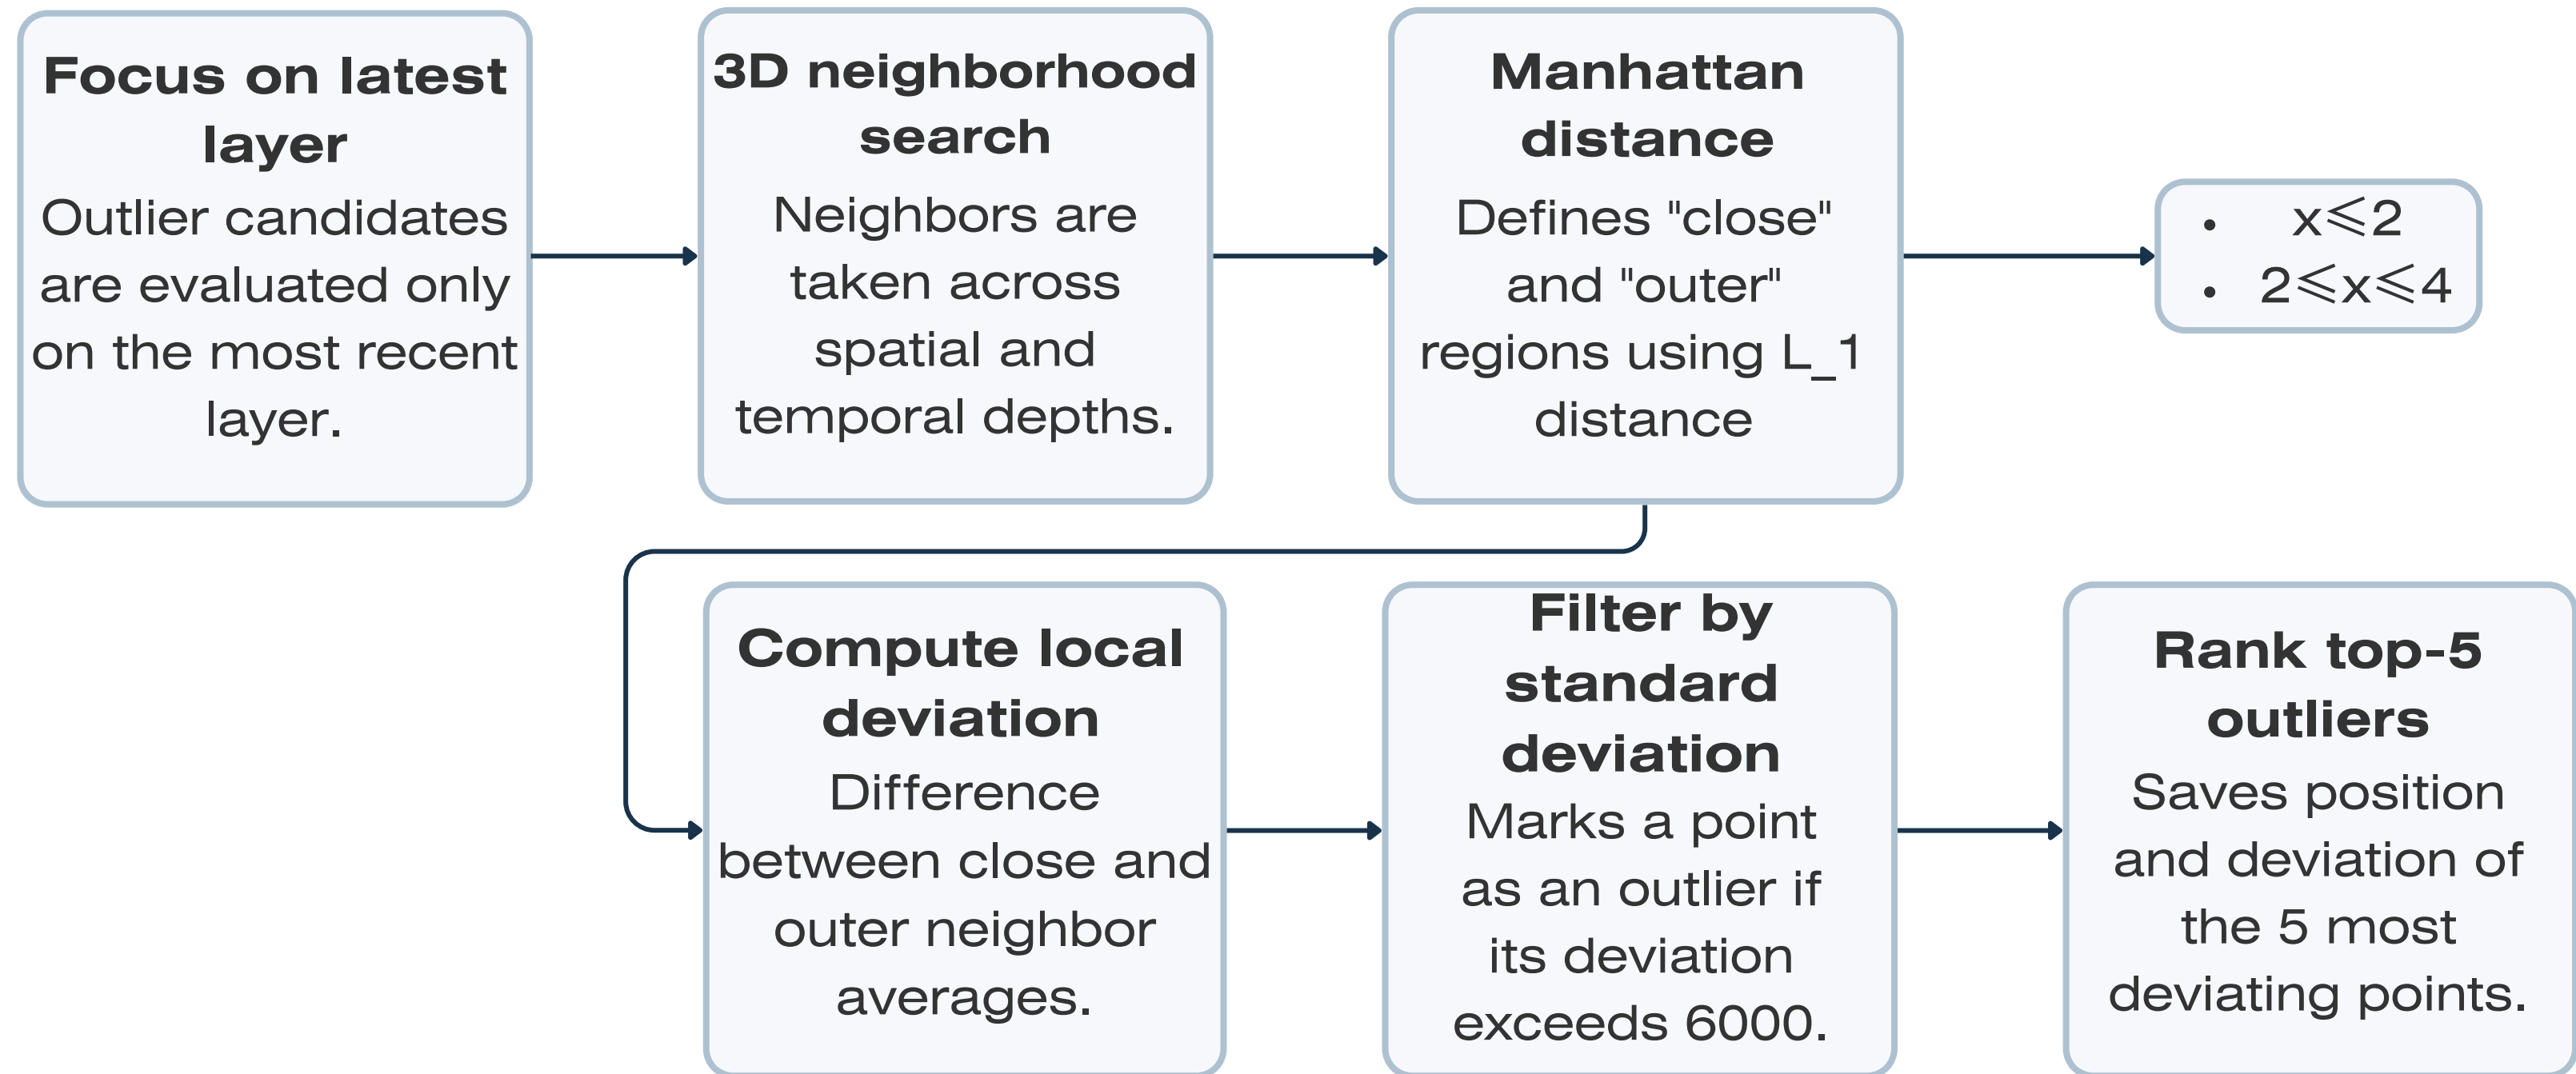
Query 2 - Flink (Windowing)



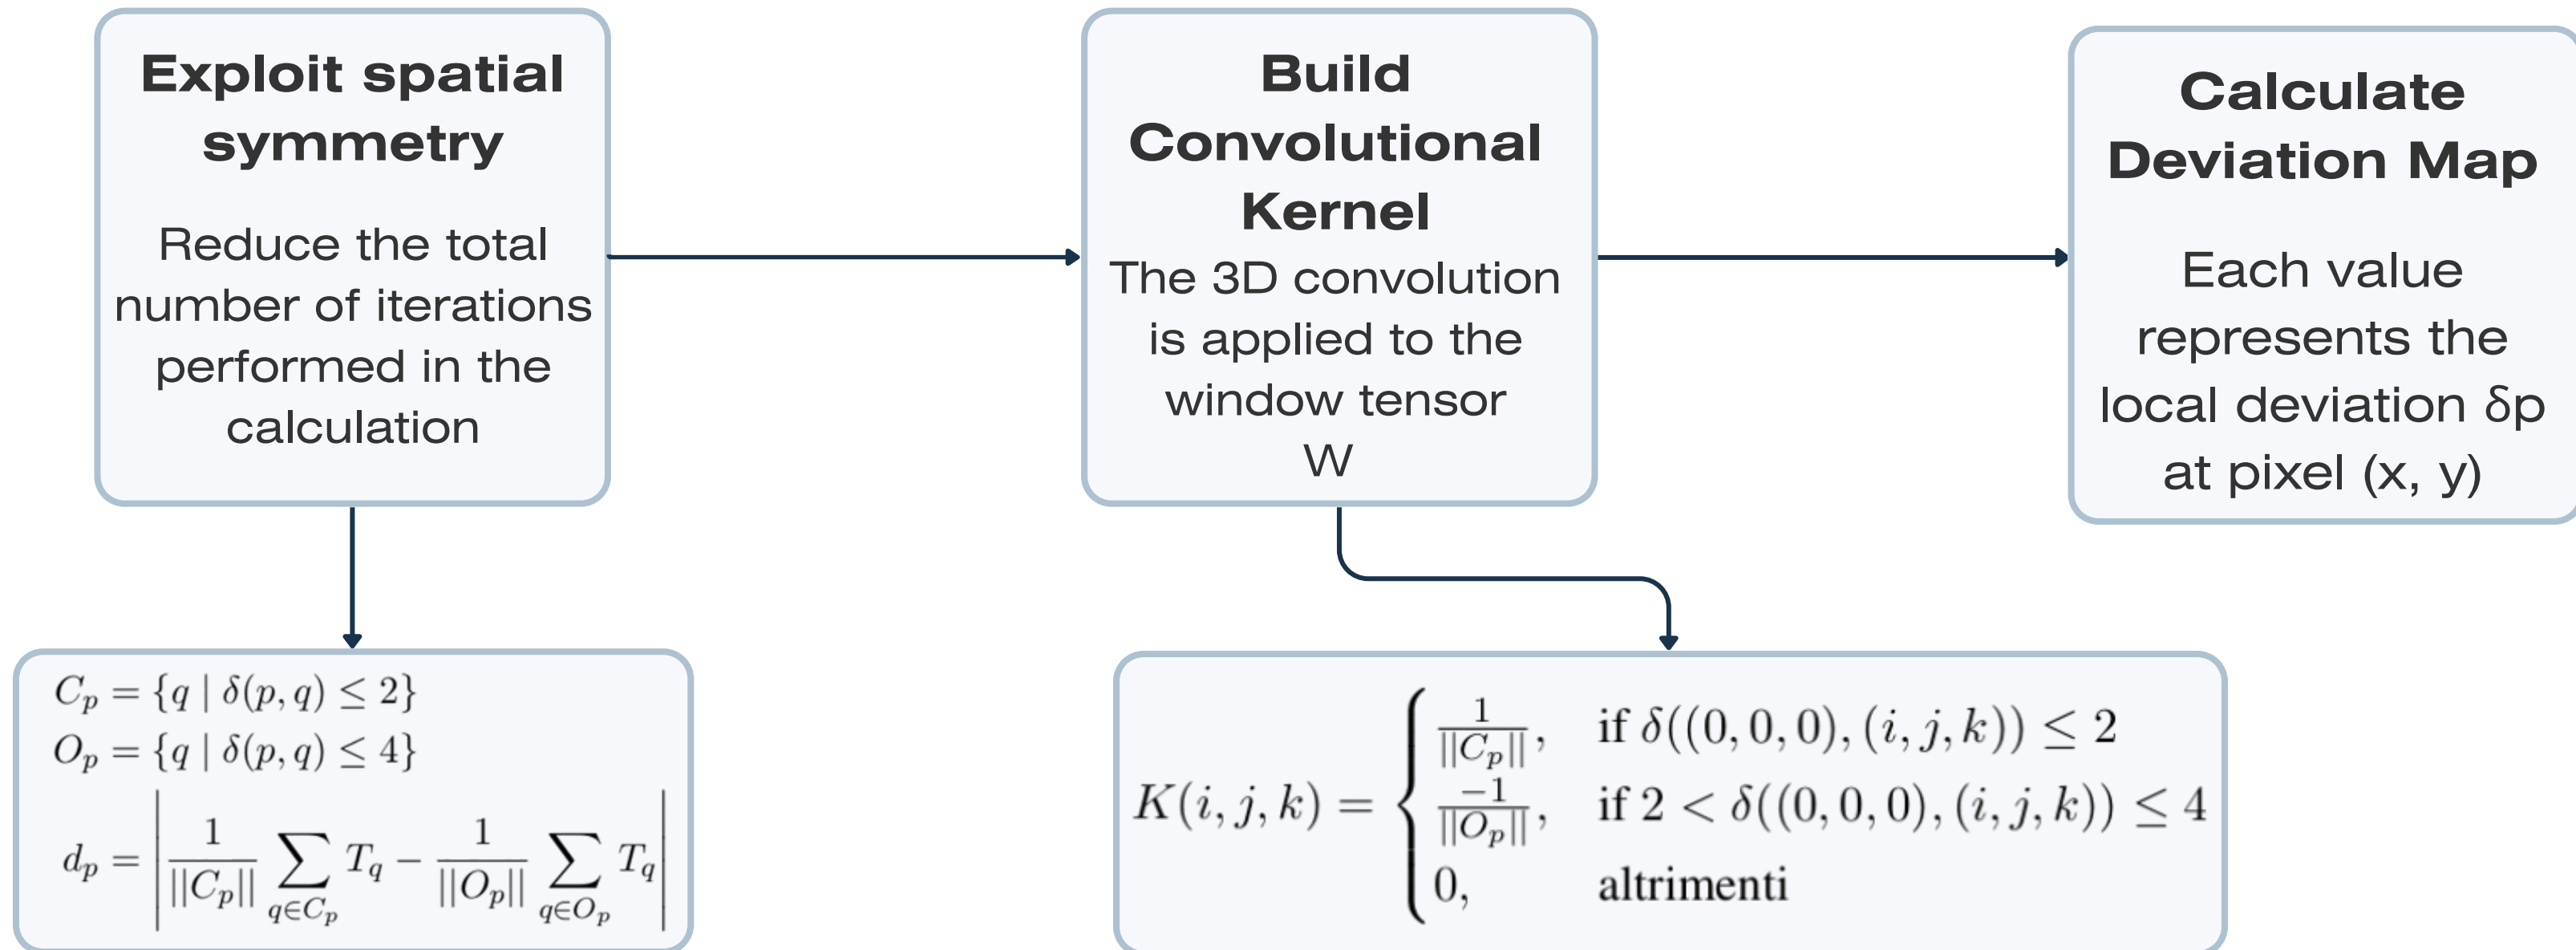
Query 2 - Flink (Outlier analysis)



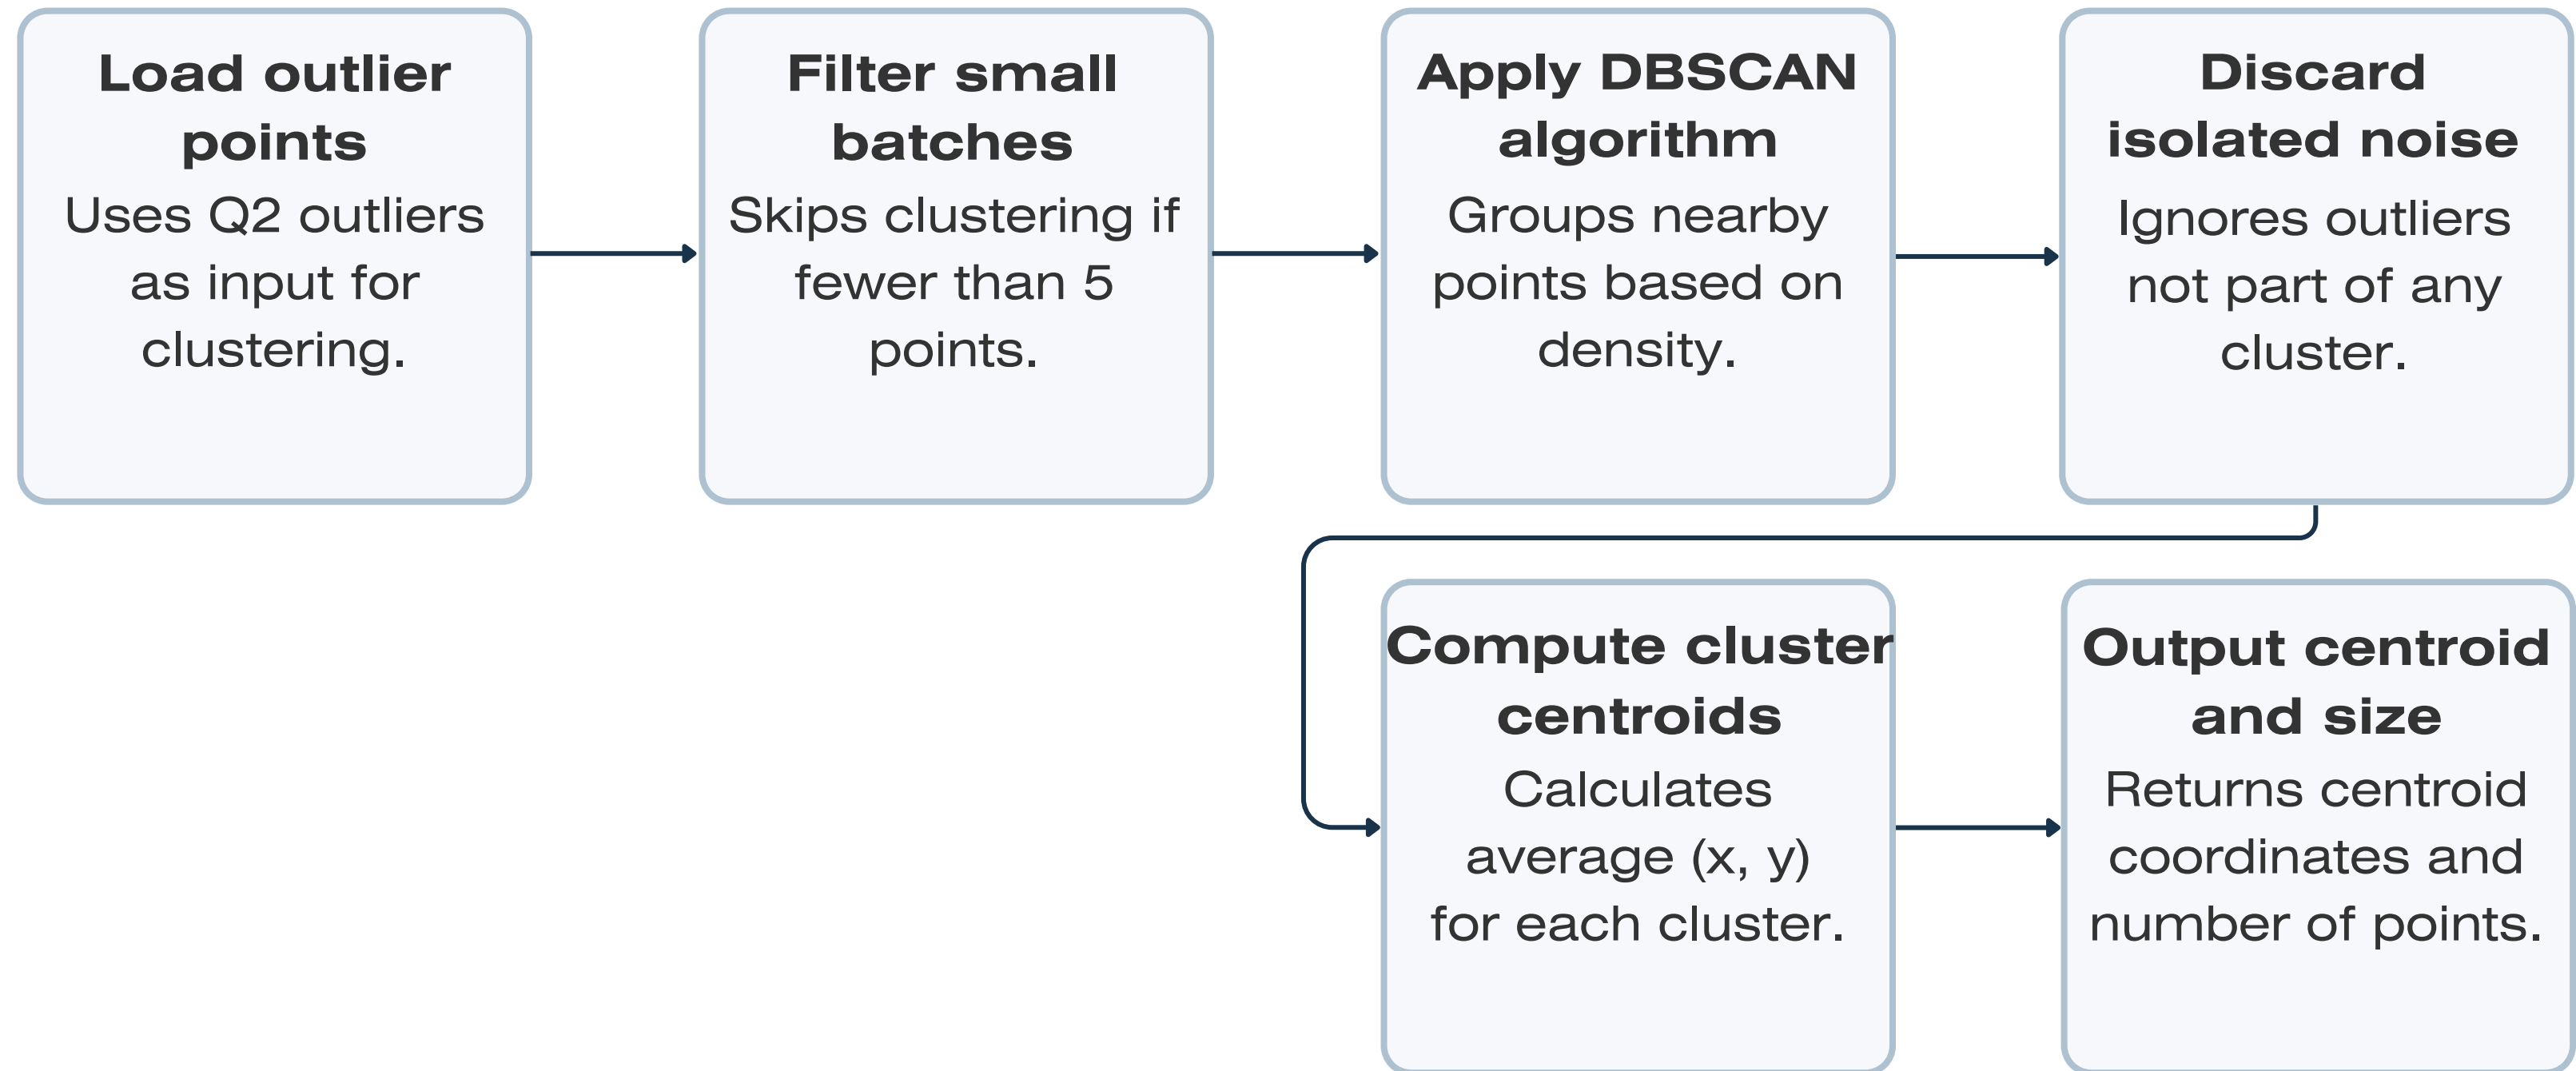
Query 2 - Flink (Outlier analysis)



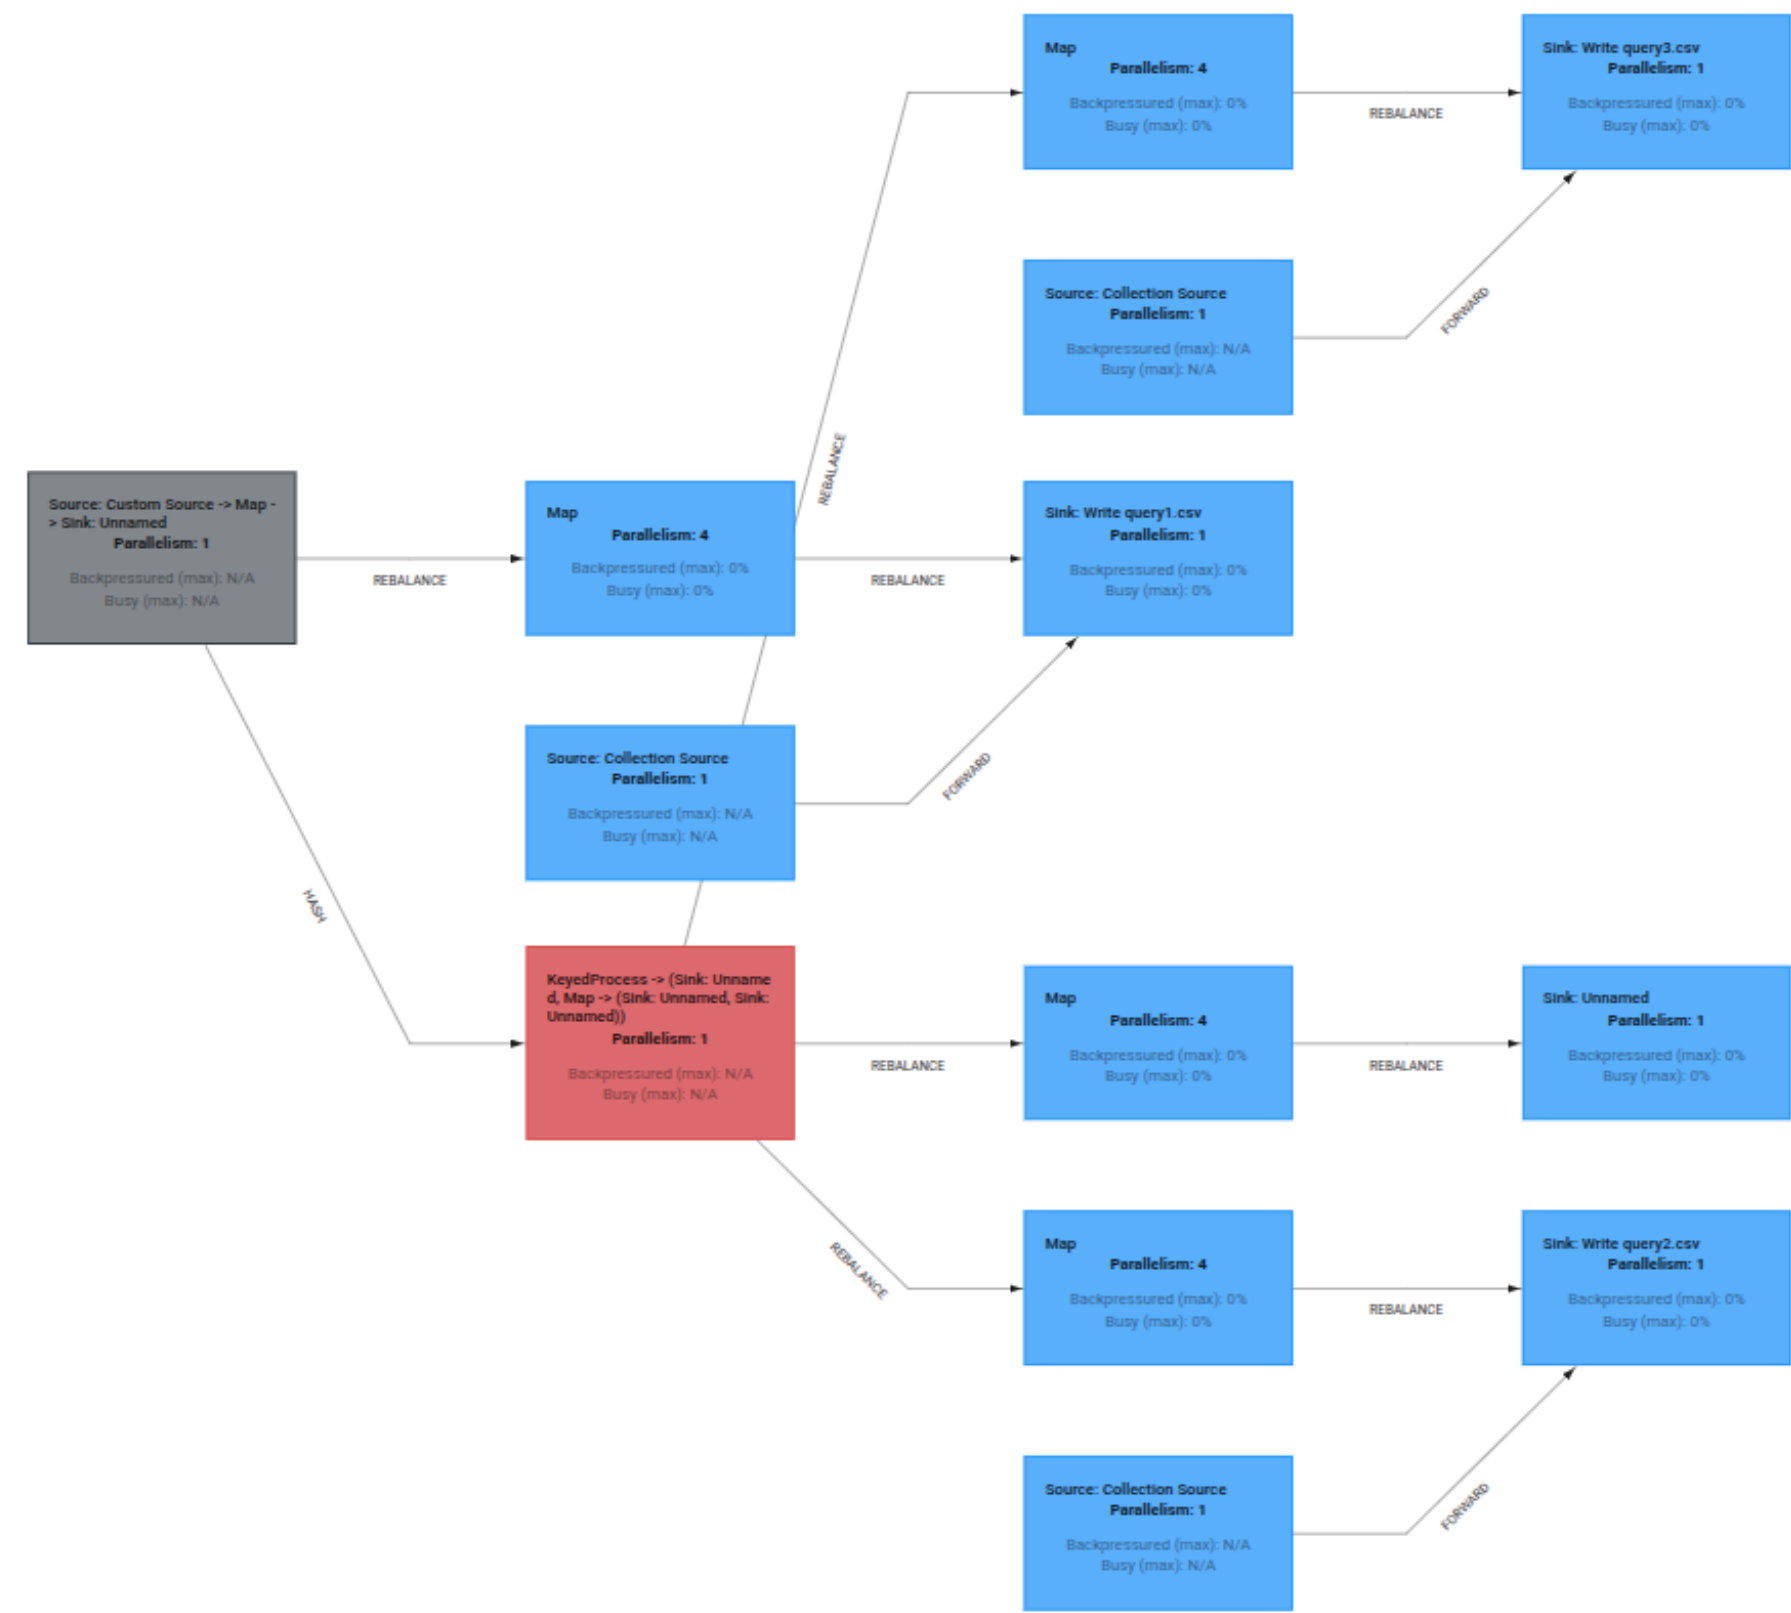
Query 2 Optimized - Flink





Query 3 - Flink



Flink - DAG



Results - Flink Version

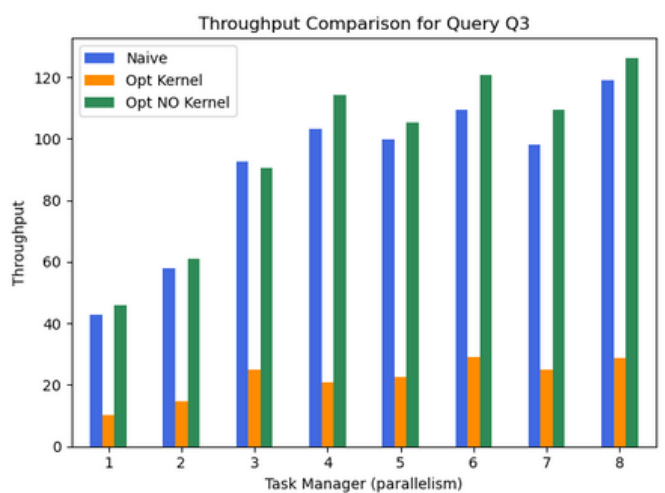



Results - Flink Version




	parallelism ▼	÷	bench_id ▼	÷	throughput ▼	÷	latency_mean ▼	÷	latency_max ▼	÷
1		1	01JZFC54SBDP2TKRR47KN65S7Y		42.80		45ms429µs638ns		226ms26µs220ns	
2		2	01JZFCVQT4BWXMNRTN530P1XZ		57.92		150ms608µs245ns		359ms985µs929ns	
3		3	01JZFCXNESZPABDJ22EHNMRWCQ		92.33		156ms928µs791ns		440ms297µs384ns	
4		4	01JZFCYXEVMKKQ741WAB1R3VA		102.62		163ms190µs819ns		490ms156µs78ns	
5		5	01JZFD01VNMFA2YKASDVM6TNJS		99.14		196ms580µs312ns		581ms91µs858ns	
6		6	01JZFD15ZDC99MQ7A0N3PBZG8Z		109.31		203ms833µs977ns		713ms358µs40ns	
7		7	01JZFD28978RGHDQZ0YHWP2B38		97.81		236ms257µs818ns		793ms698µs202ns	
8		8	01JZFD3EVTAKB0FF1CJYCZDMD4		118.48		205ms153µs826ns		673ms430µs230ns	
9		1	01JZMS1NPQ9D87D4TKY67ZV84X		10.11		195ms810µs931ns		336ms526µs330ns	
10		2	01JZMSCK3F70WW0WDNW48VAEN2		14.55		574ms802µs902ns		1s353ms600µs274ns	
11		3	01JZMSM6H6M1XQMSZKPKV6BR8C		25.07		497ms869µs679ns		1s267ms758µs257ns	
12		4	01JZMSRM9BY0REQE45Q787FMYX		20.71		500ms178µs369ns		1s293ms715µs701ns	
13		5	01JZMSXZC3FBCPSJ62F63PF54T		22.59		732ms853µs132ns		2s171ms416µs910ns	
14		6	01JZMT2WSGCPHPBC368A0ZJJ6N		29.17		568ms810µs971ns		1s464ms497µs679ns	
15		7	01JZMT6PSKPZRM0V0ZB9G5GZJD		25.14		775ms841µs351ns		2s381ms154µs852ns	
16		8	01JZMTB4JMQW3NC7Z9NPC4ZESX		29.01		543ms312µs256ns		2s418ms807µs851ns	
17		1	01JZFD4ED571JZNMR6PN1G1EF2		45.96		42ms171µs852ns		118ms387µs586ns	
18		2	01JZFD6VSTJ2TBRYQJZ9Z84SHP		60.96		158ms307µs79ns		537ms606µs234ns	
19		3	01JZFD8QGMXTKKP1ZF7E3H1WYQ		90.22		159ms555µs576ns		447ms823µs150ns	
20		4	01JZFD9ZJ4H3VDT8DEM3Q4TKZ4		113.68		139ms63µs312ns		446ms736µs44ns	
21		5	01JZFDAZYAXMZ47PG8K2HWW0V8		104.75		190ms587µs520ns		593ms58µs31ns	
22		6	01JZFDC26DFTCC3N6QG6TMRGQCW		121.01		177ms768µs746ns		467ms259µs783ns	
23		7	01JZFDD0K5RRY86ZH10RJY5EA2		109.12		221ms23µs61ns		740ms50µs76ns	
24		8	01JZFDE2YBVMQ3SWBB1VDHJ0Z9		125.76		195ms537µs562ns		626ms853µs155ns	

Results Query - Flink Version




Throughput



Max Latency



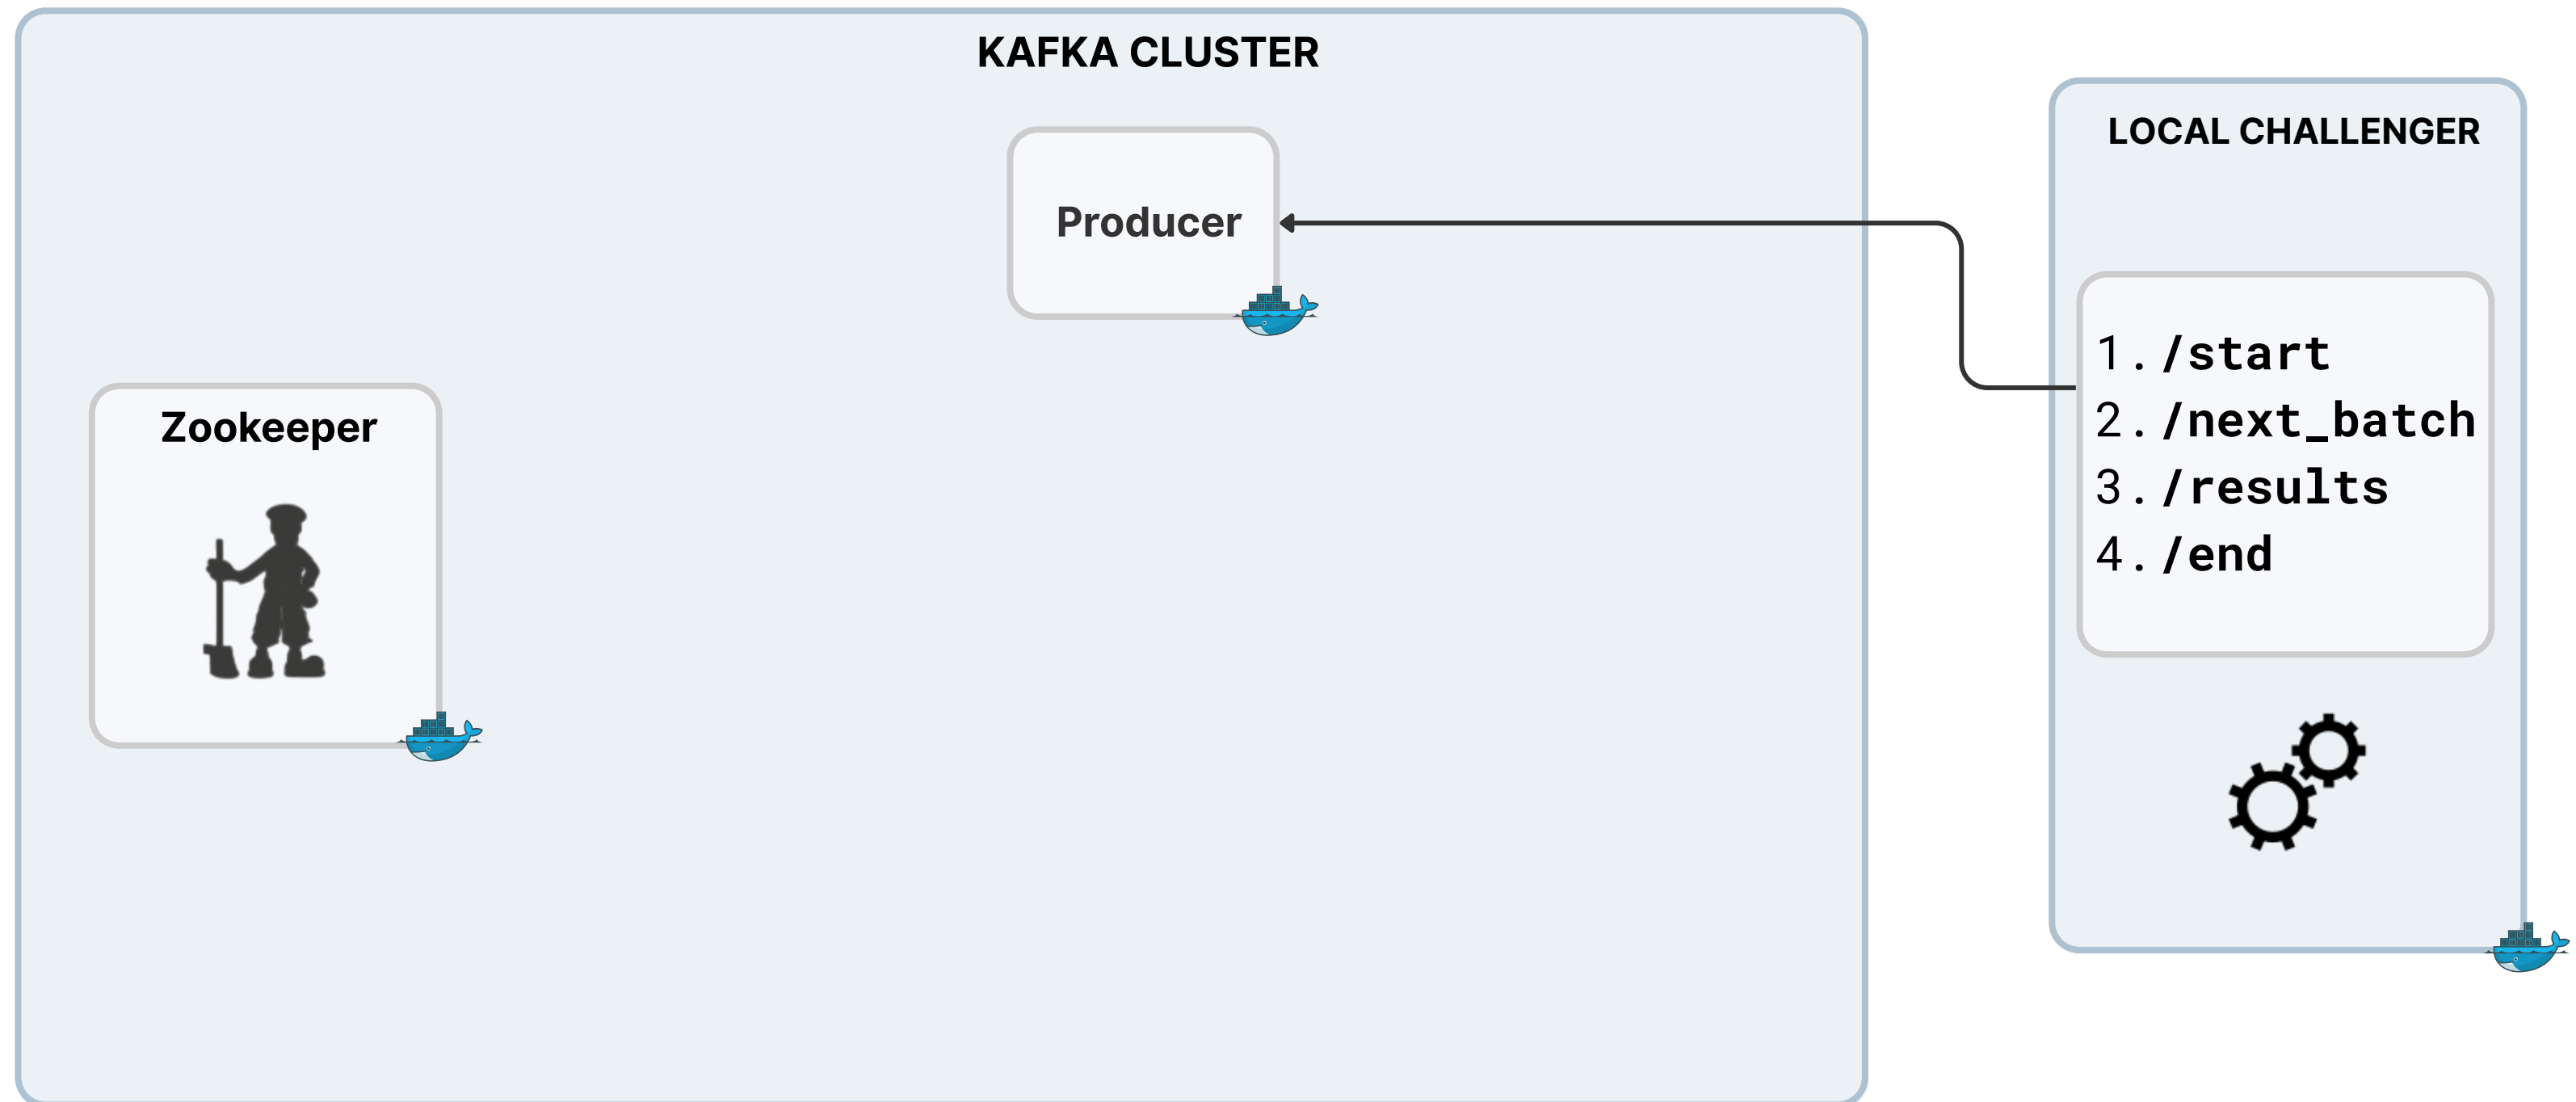
Avg Latency



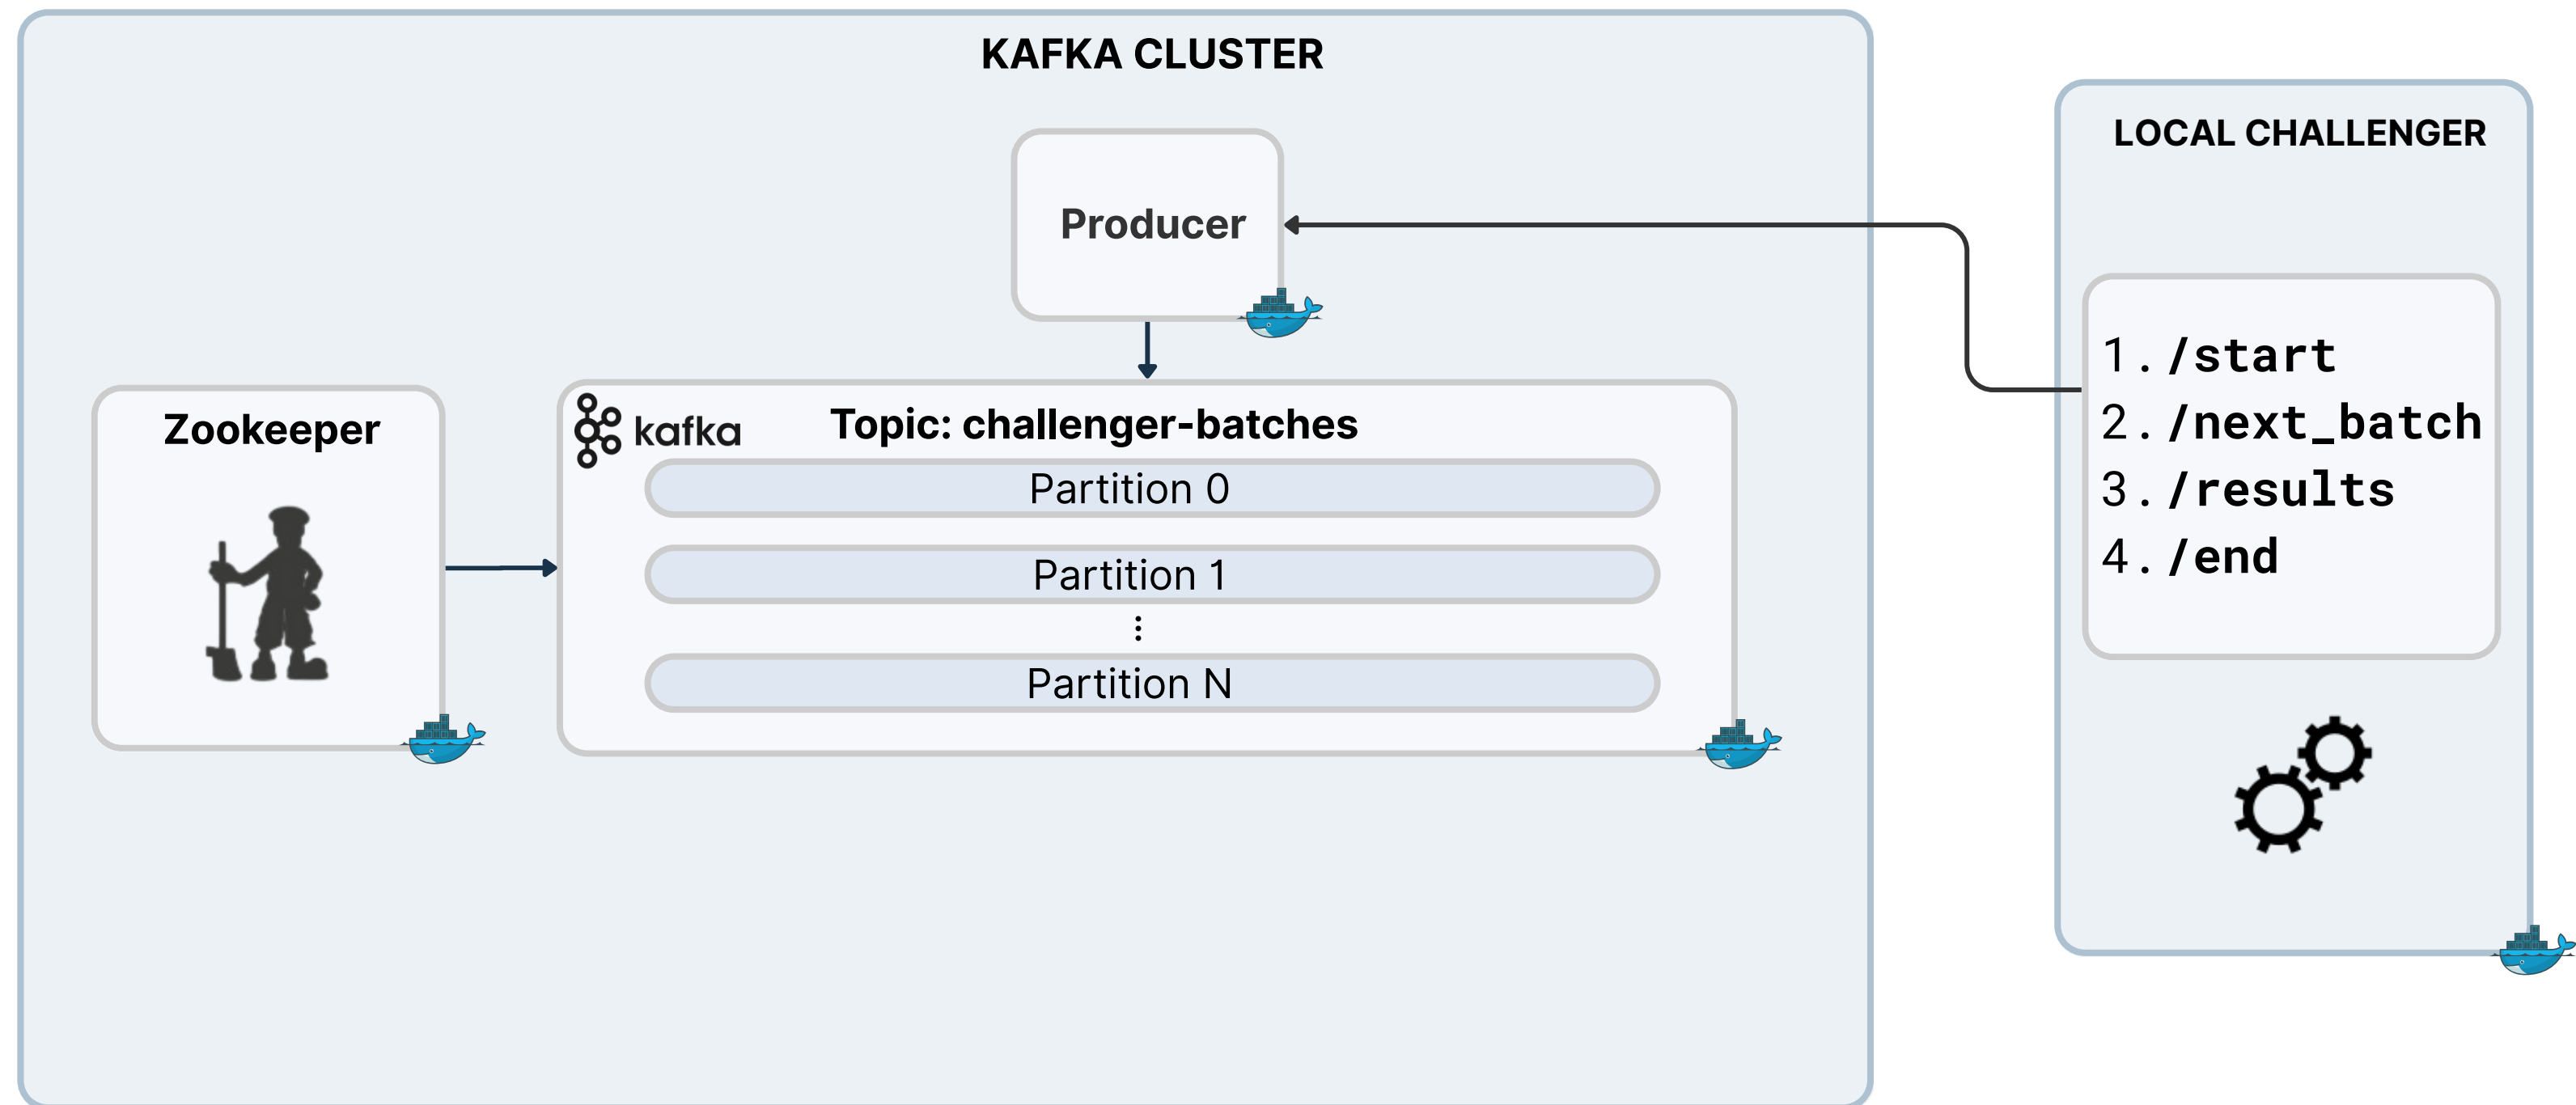
Kafka System Architecture



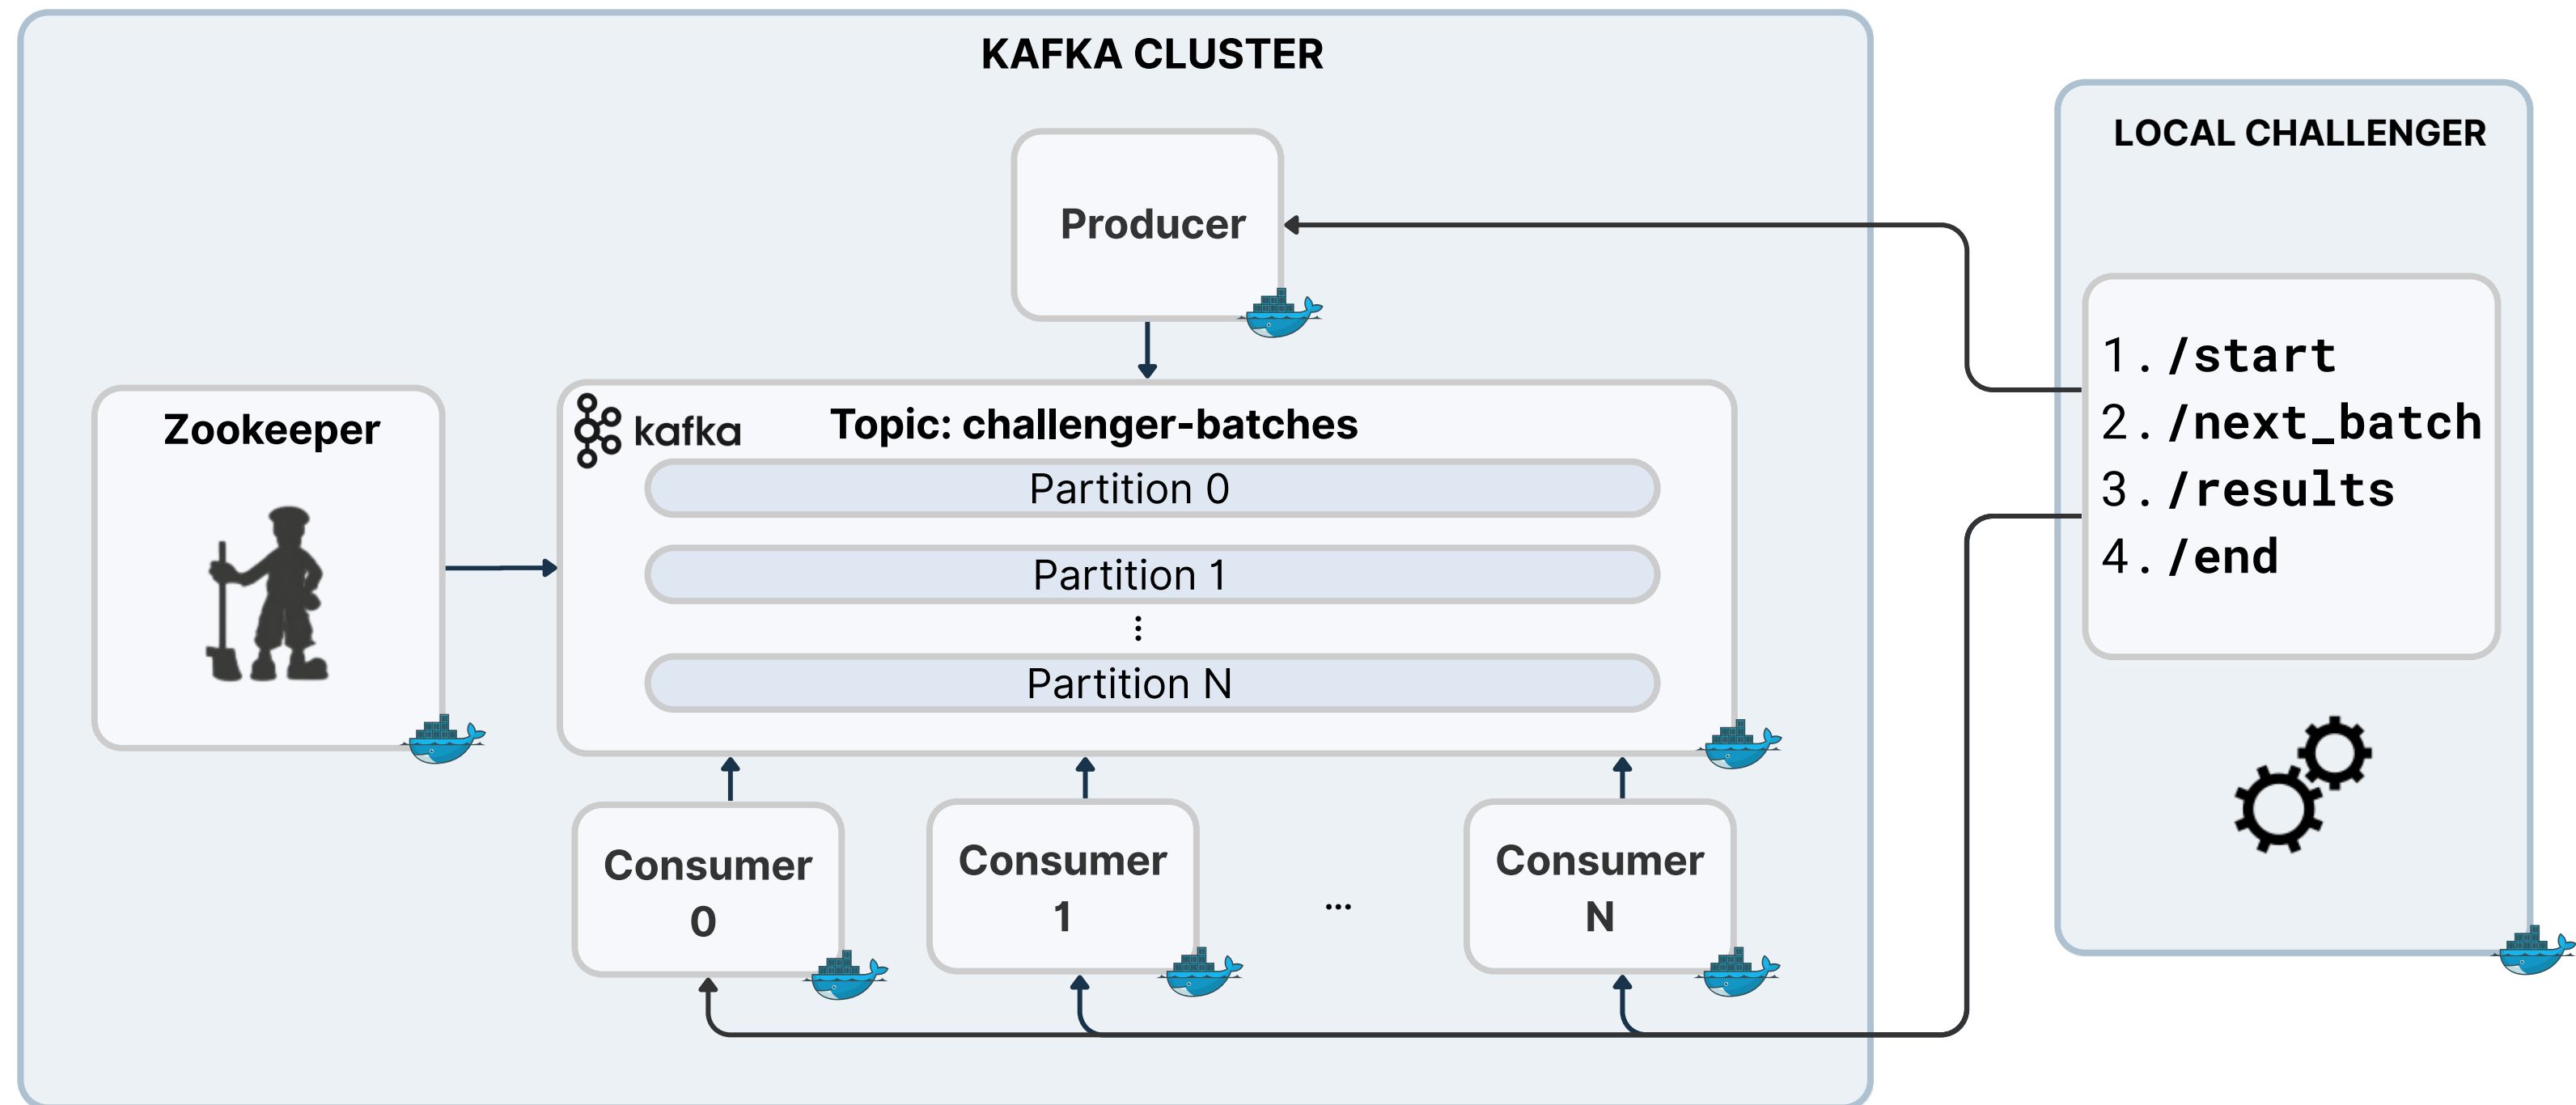
Kafka System Architecture



Kafka System Architecture

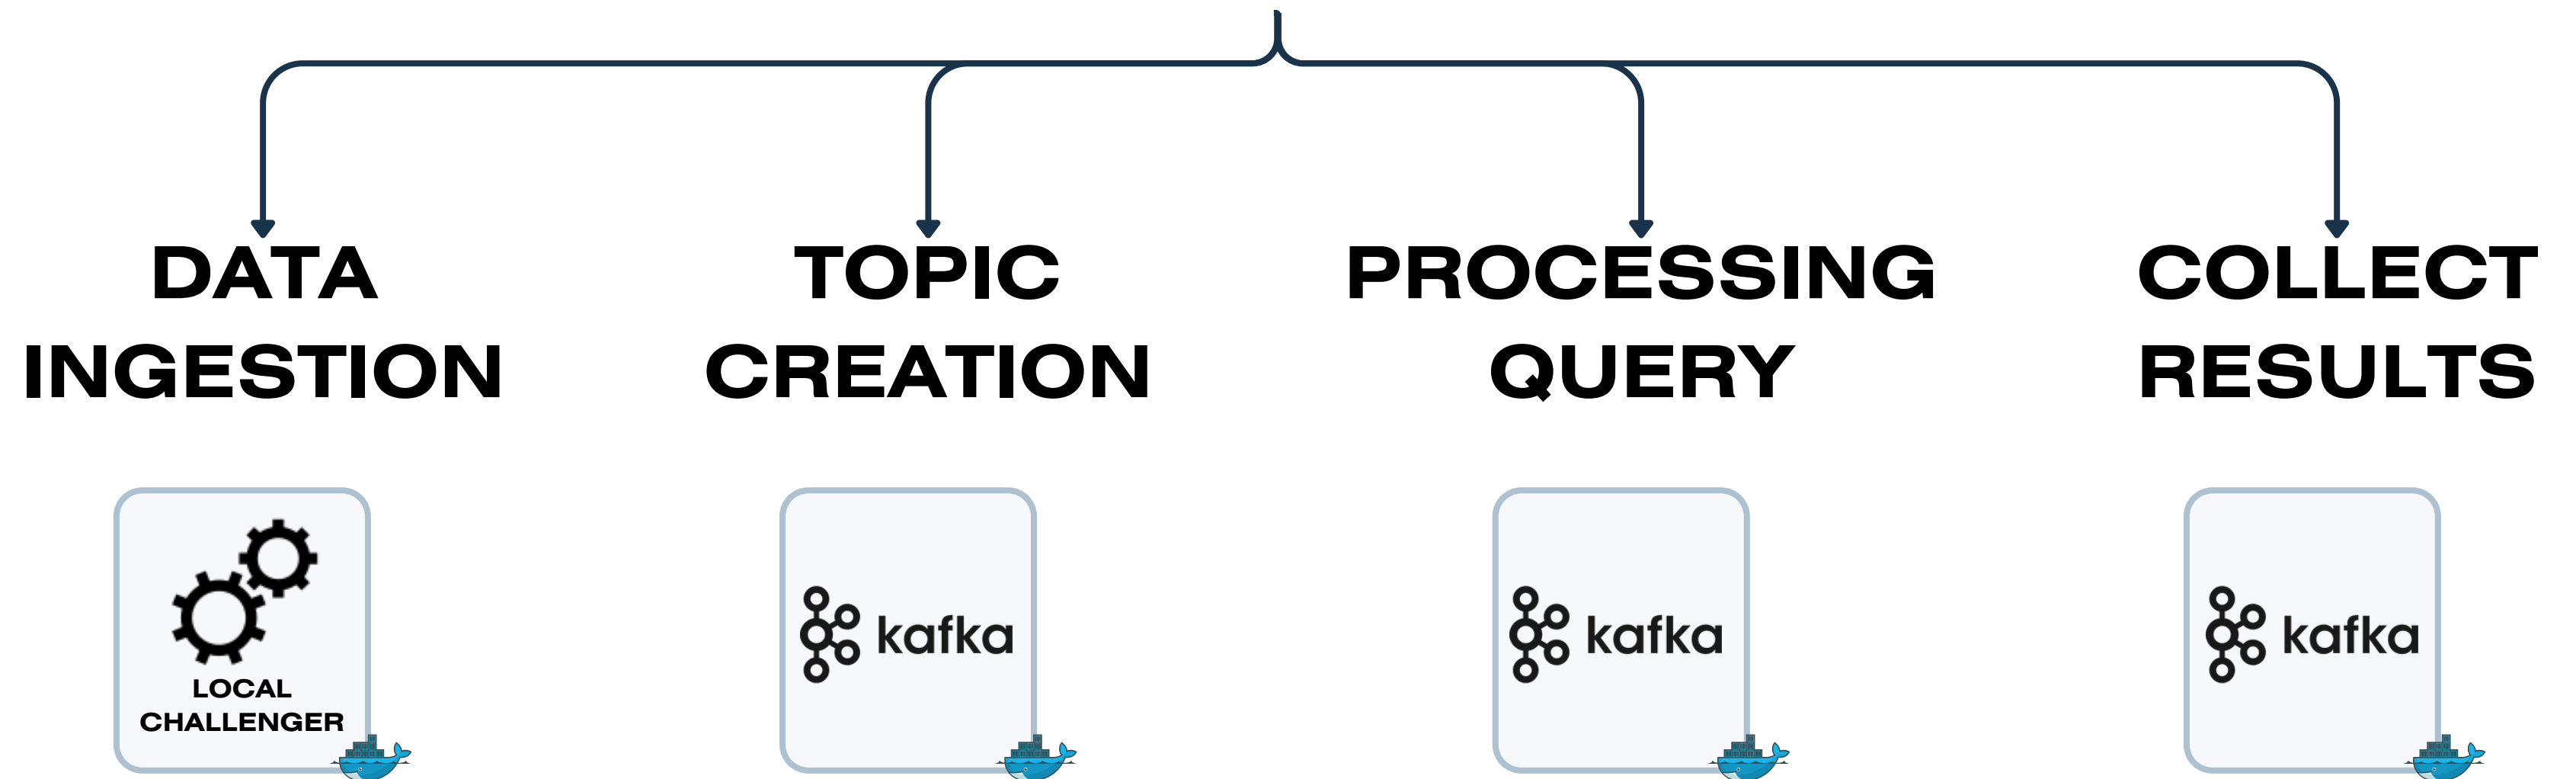


Kafka System Architecture

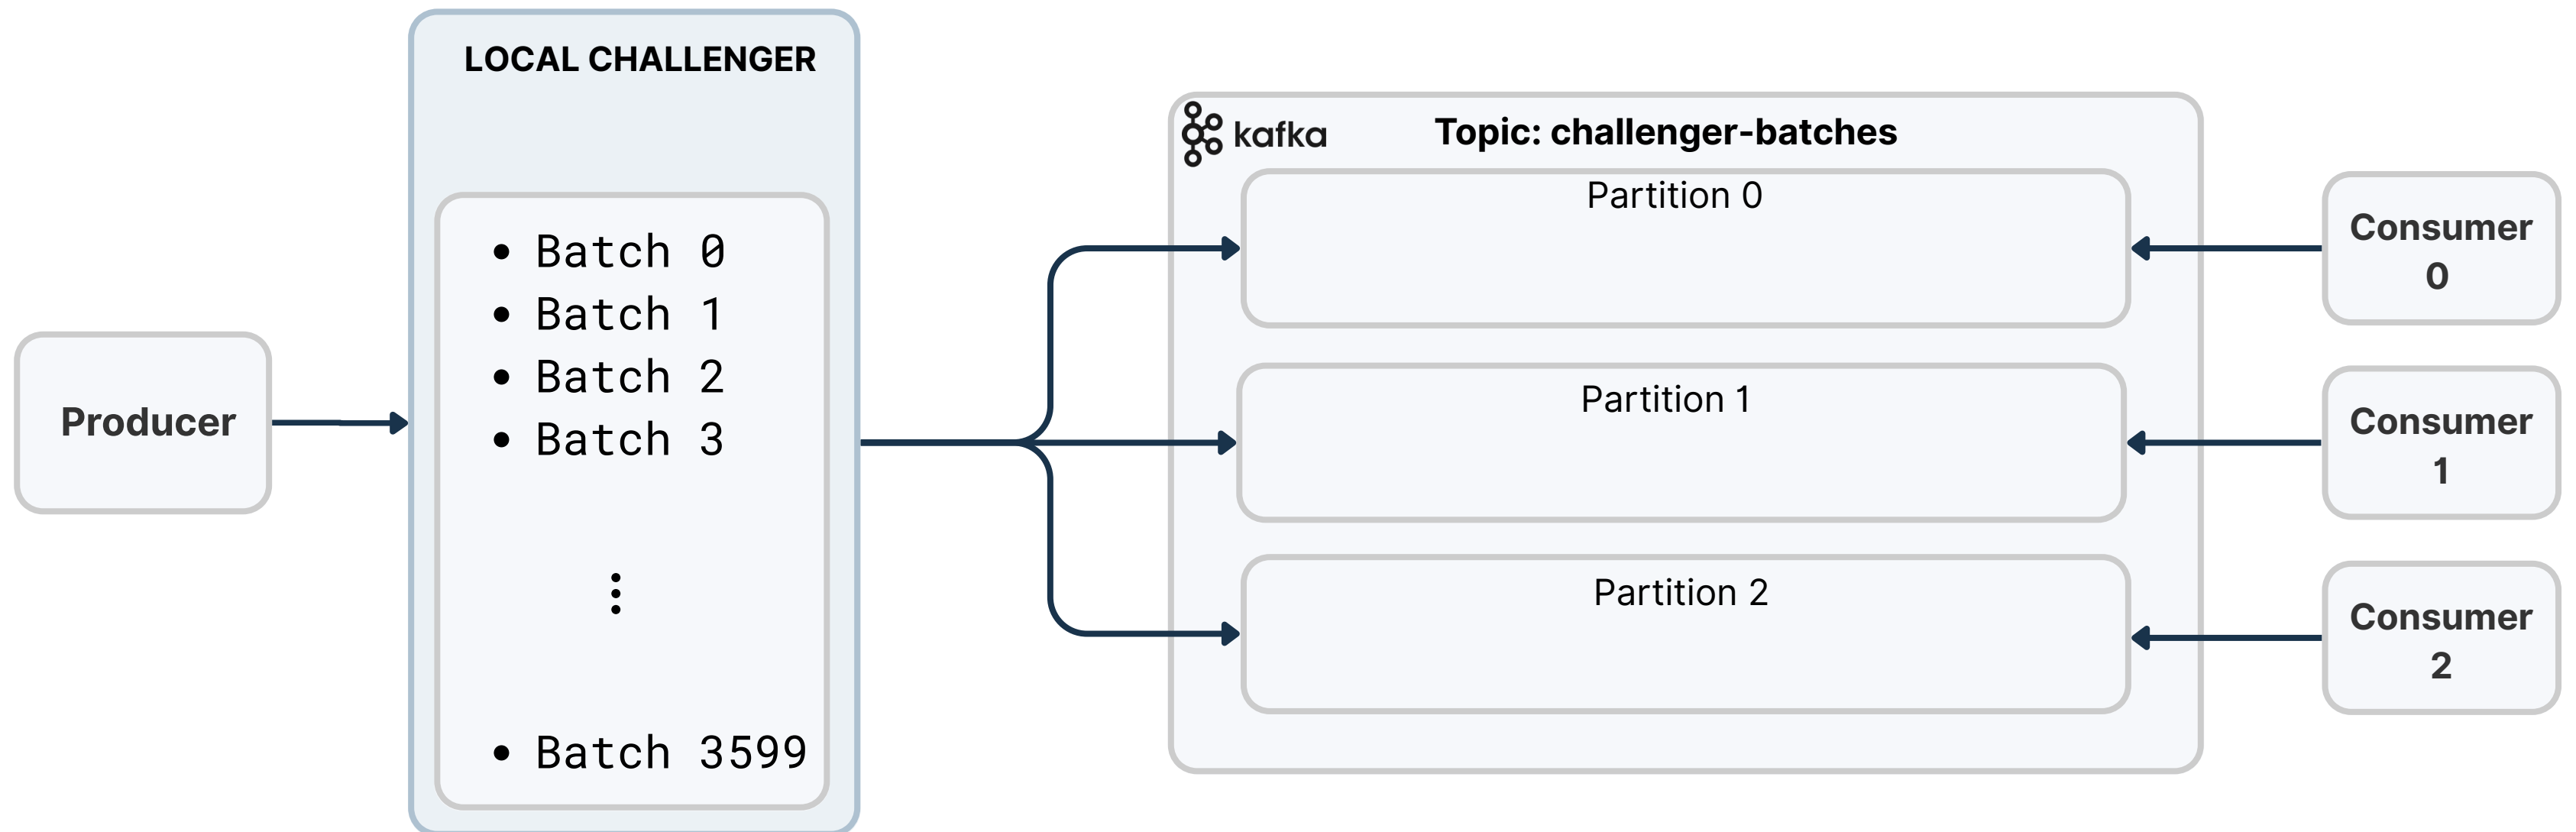


Kafka Pipeline

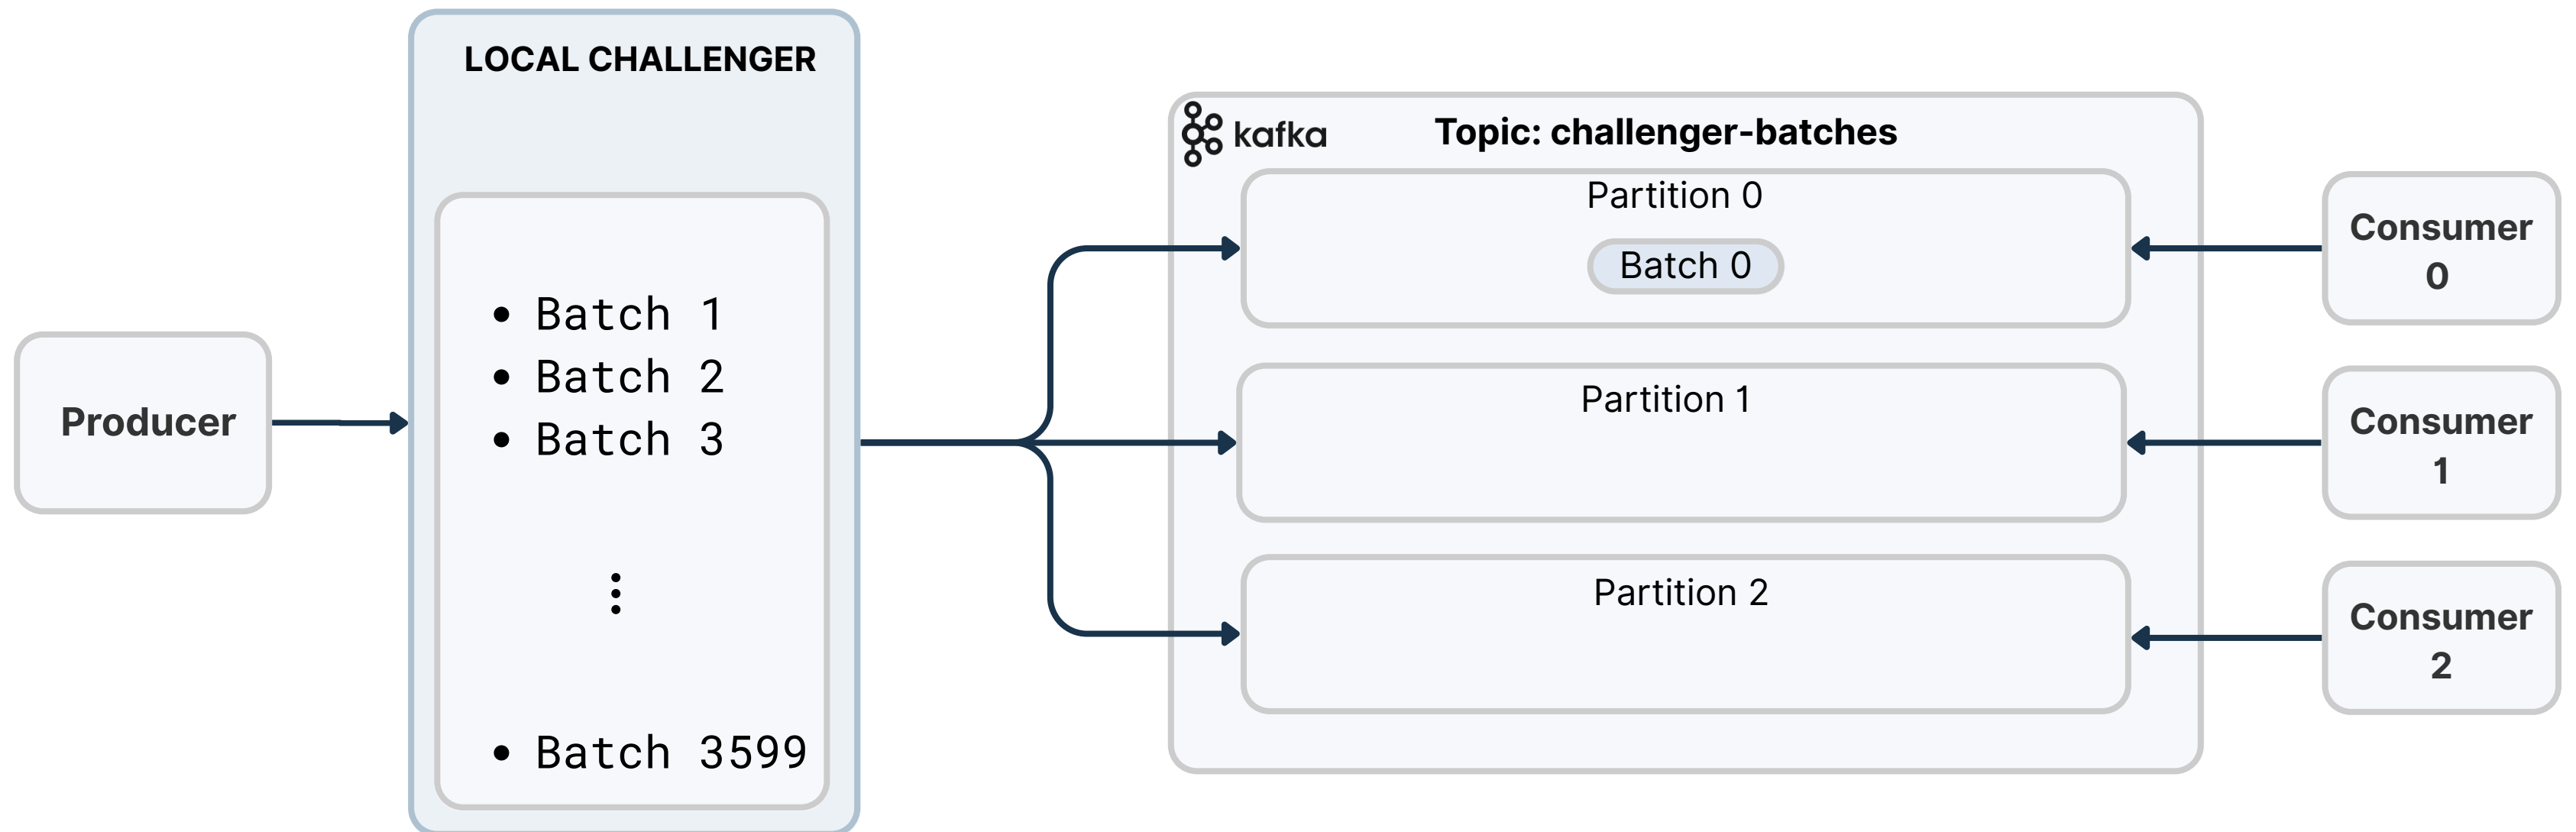
The data processing pipeline consists of the following steps:



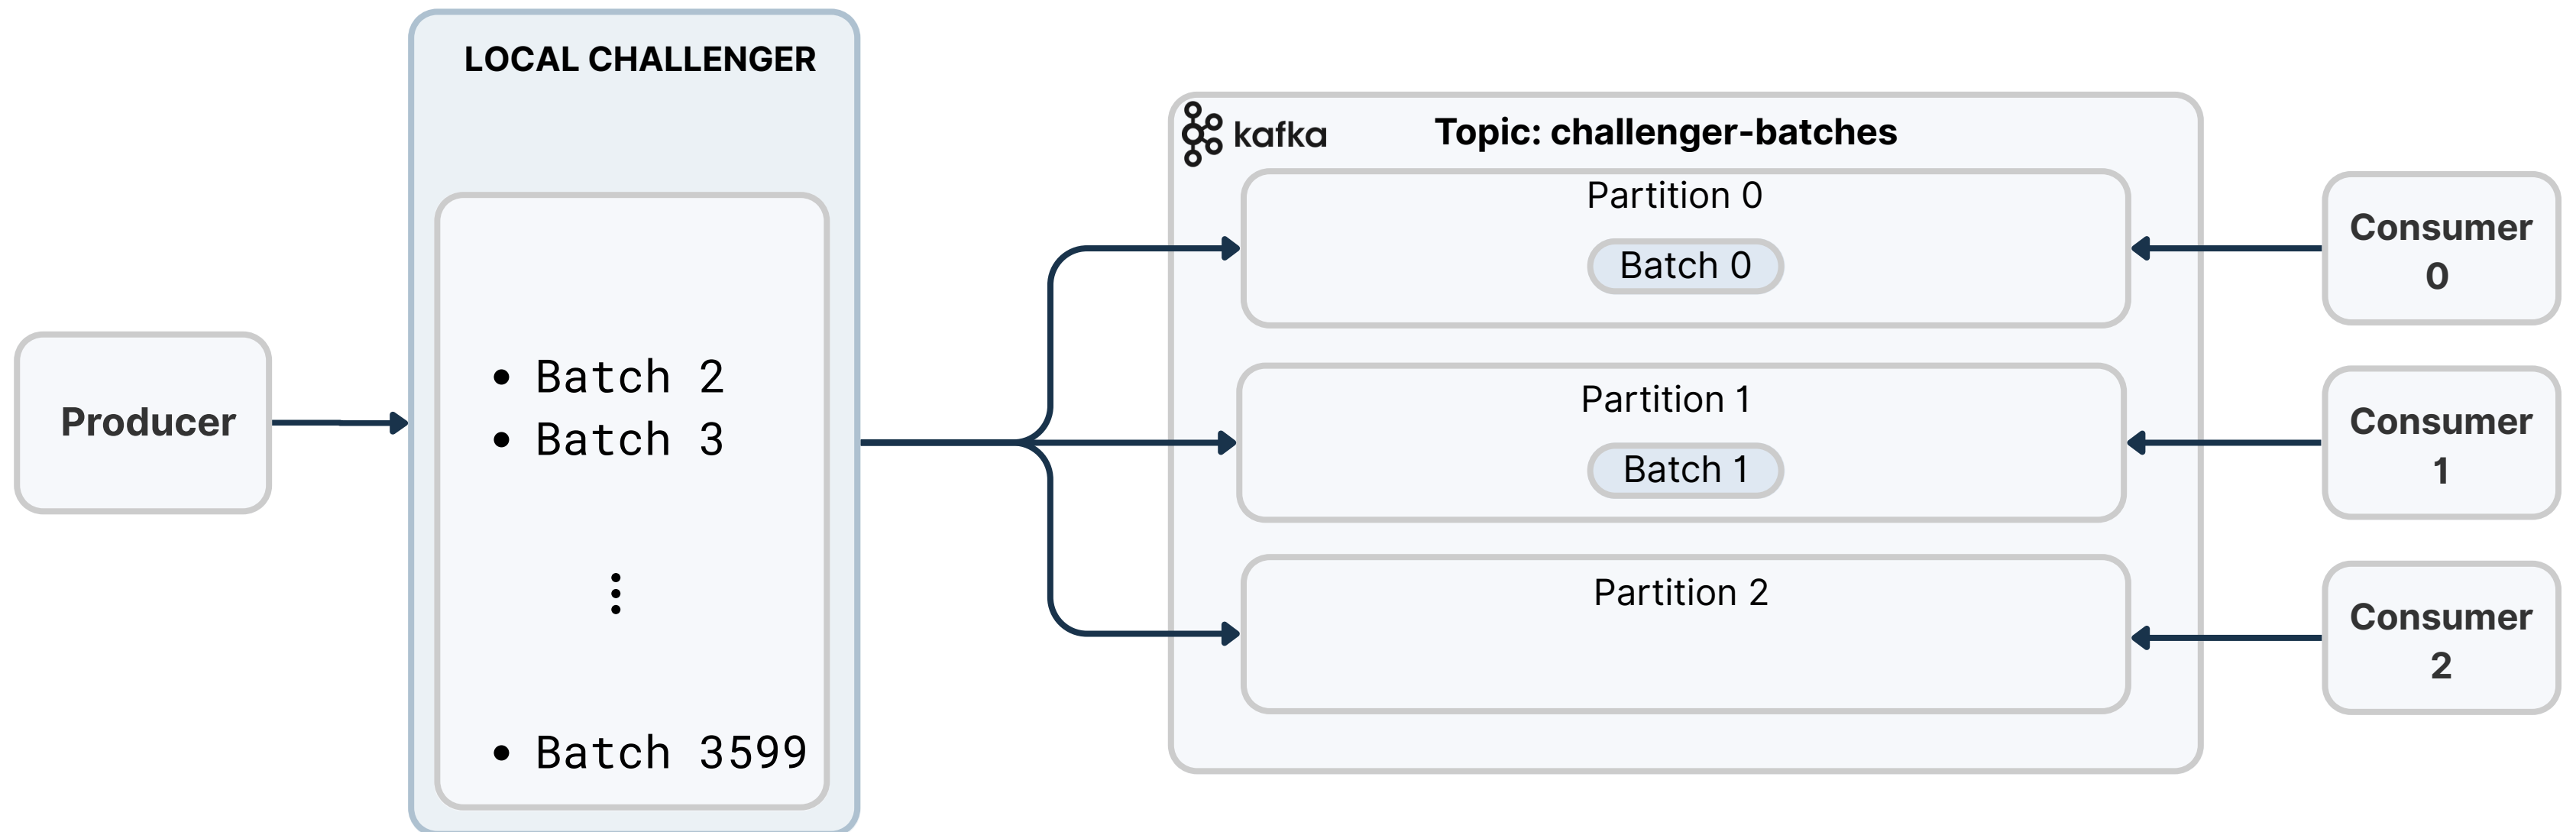
Kafka Producer



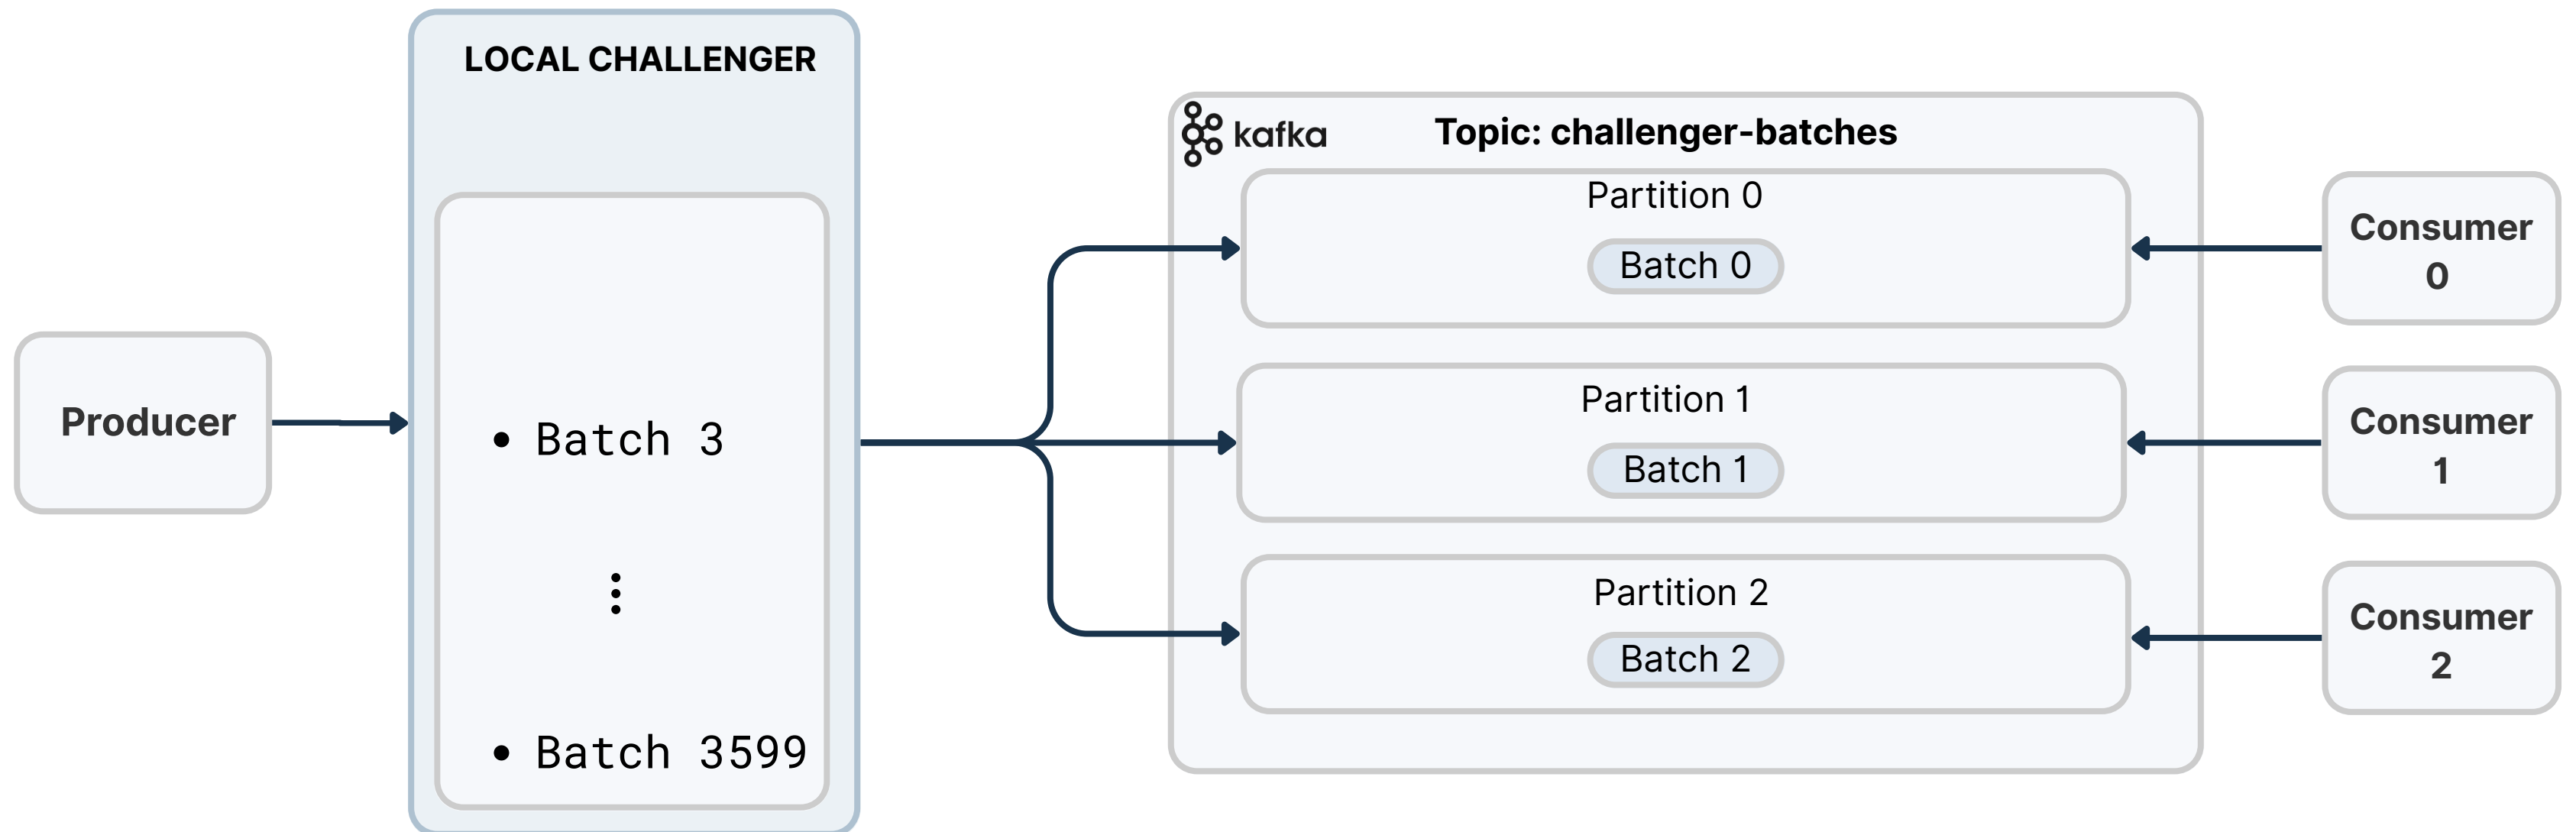
Kafka Producer



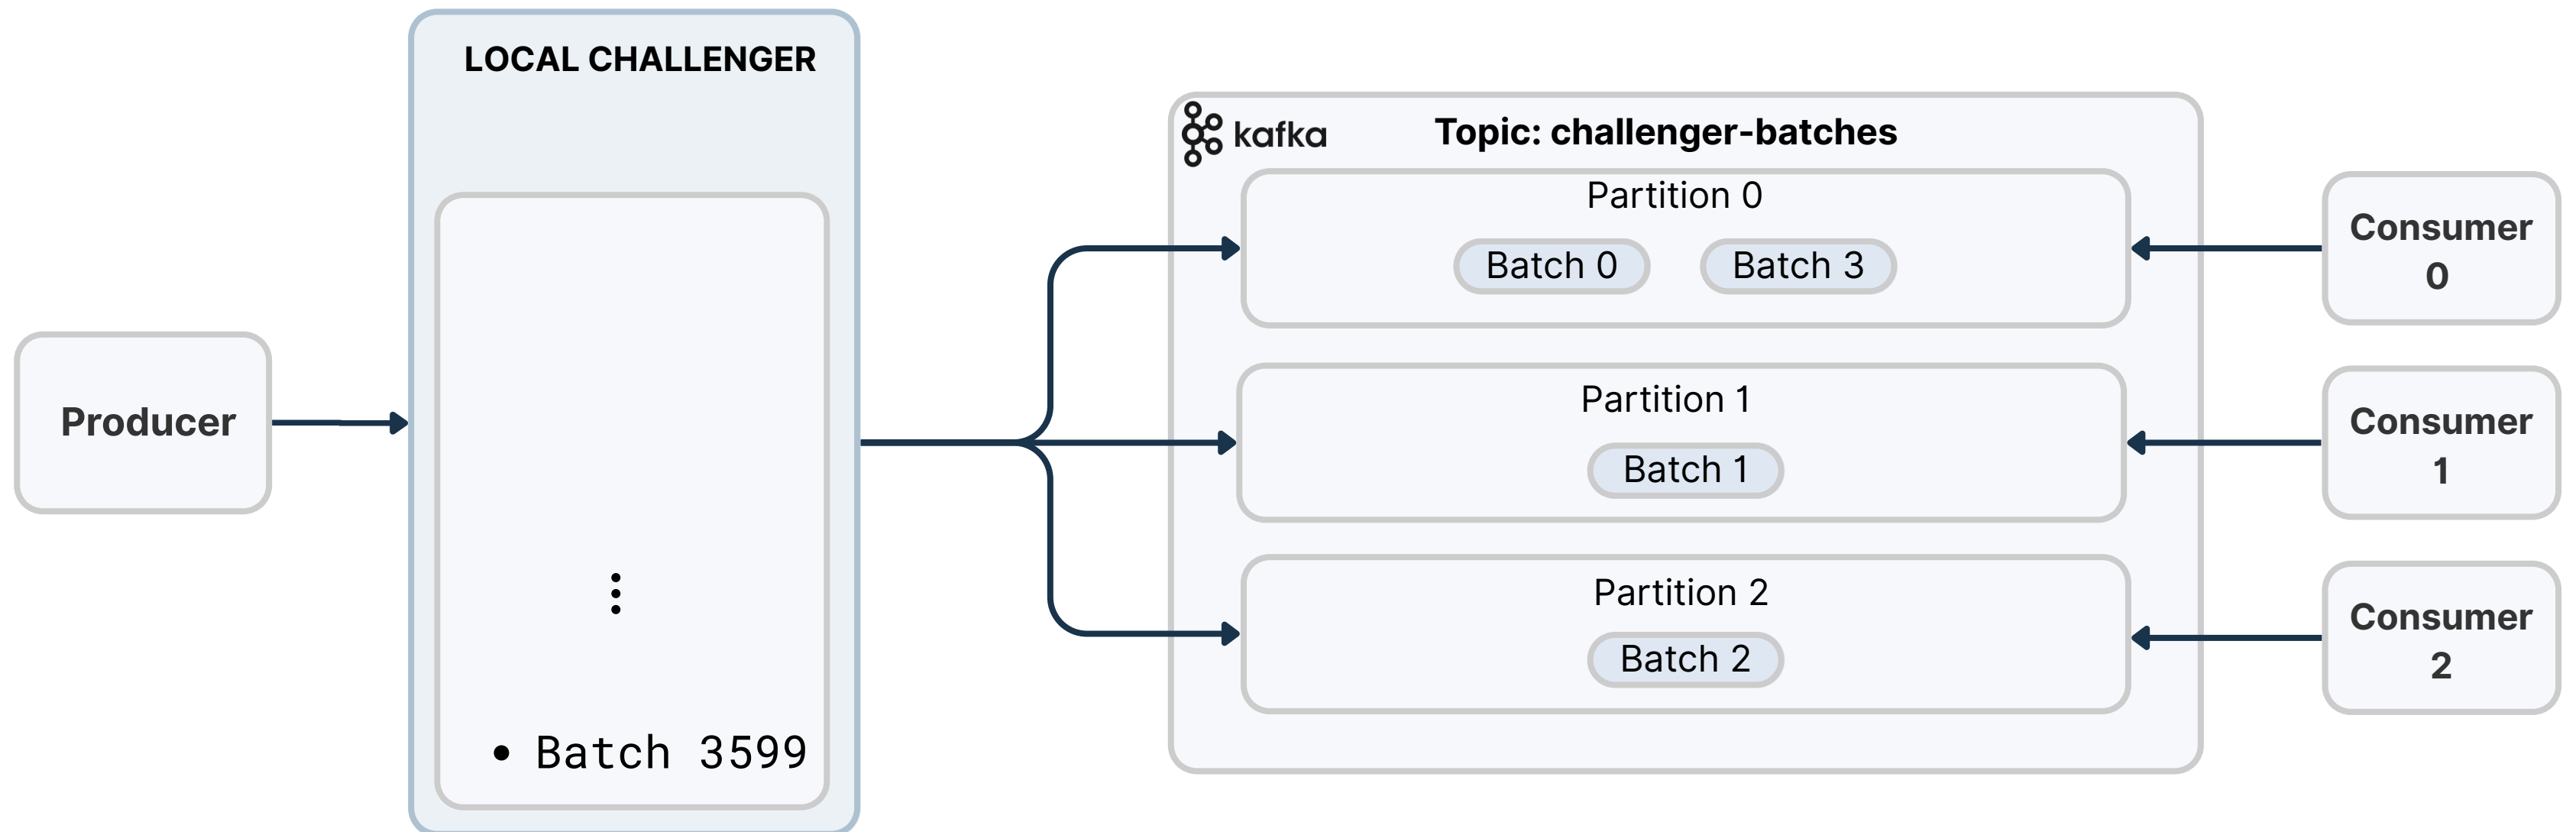
Kafka Producer



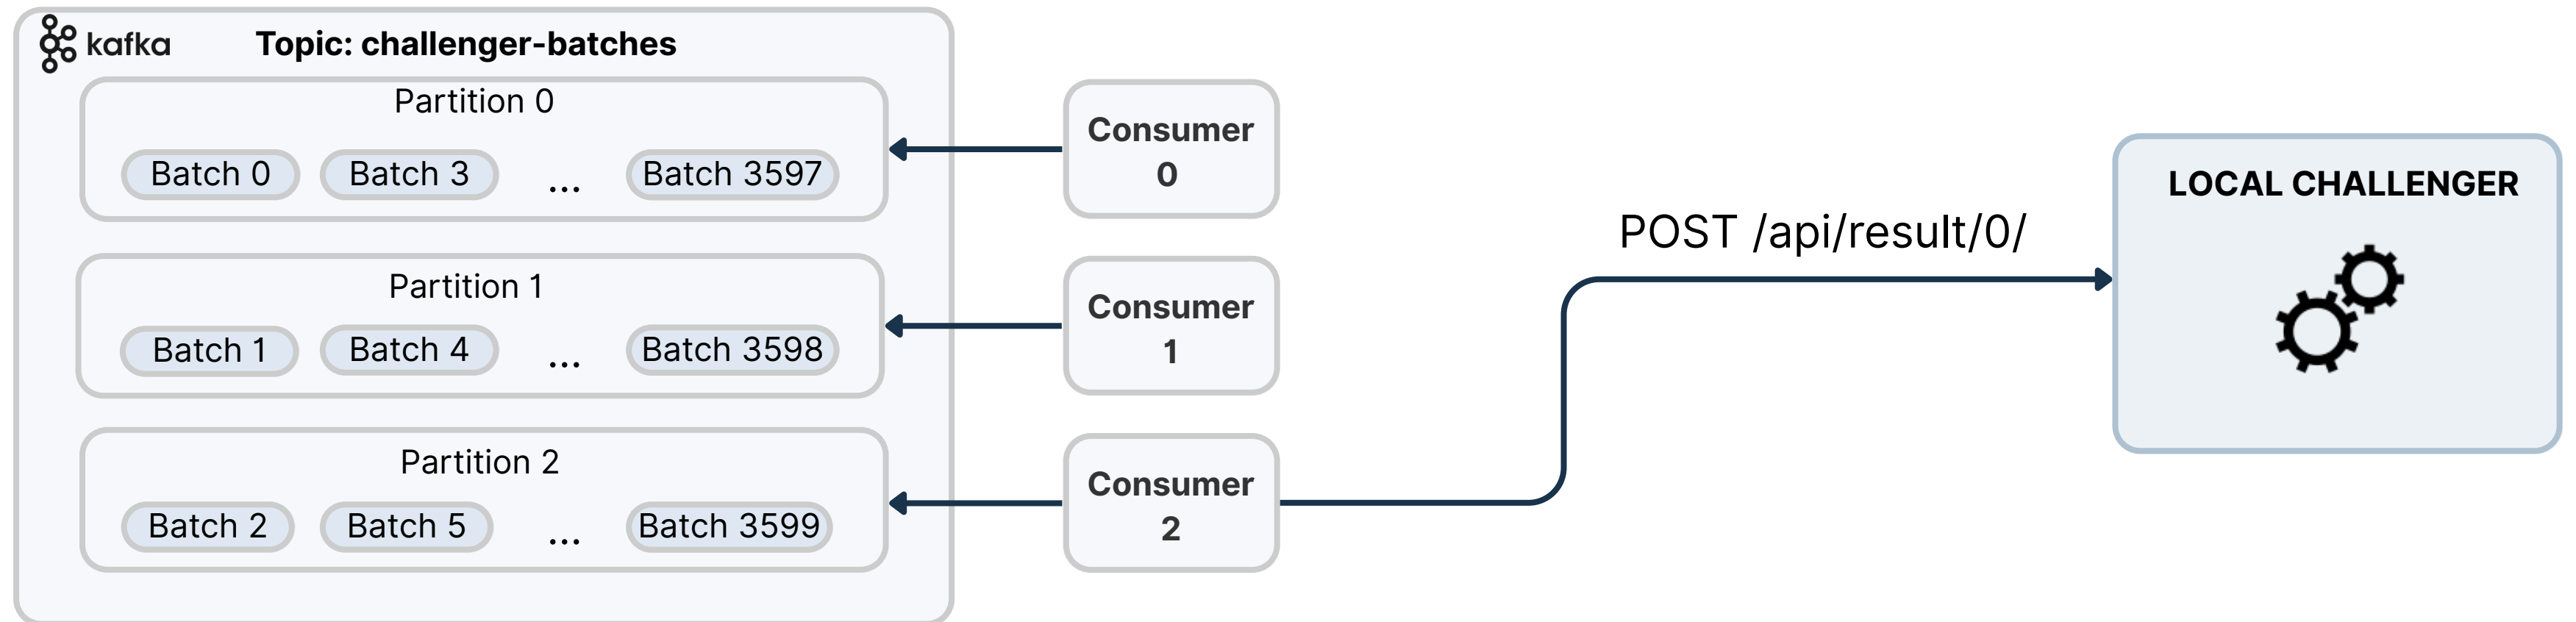
Kafka Producer



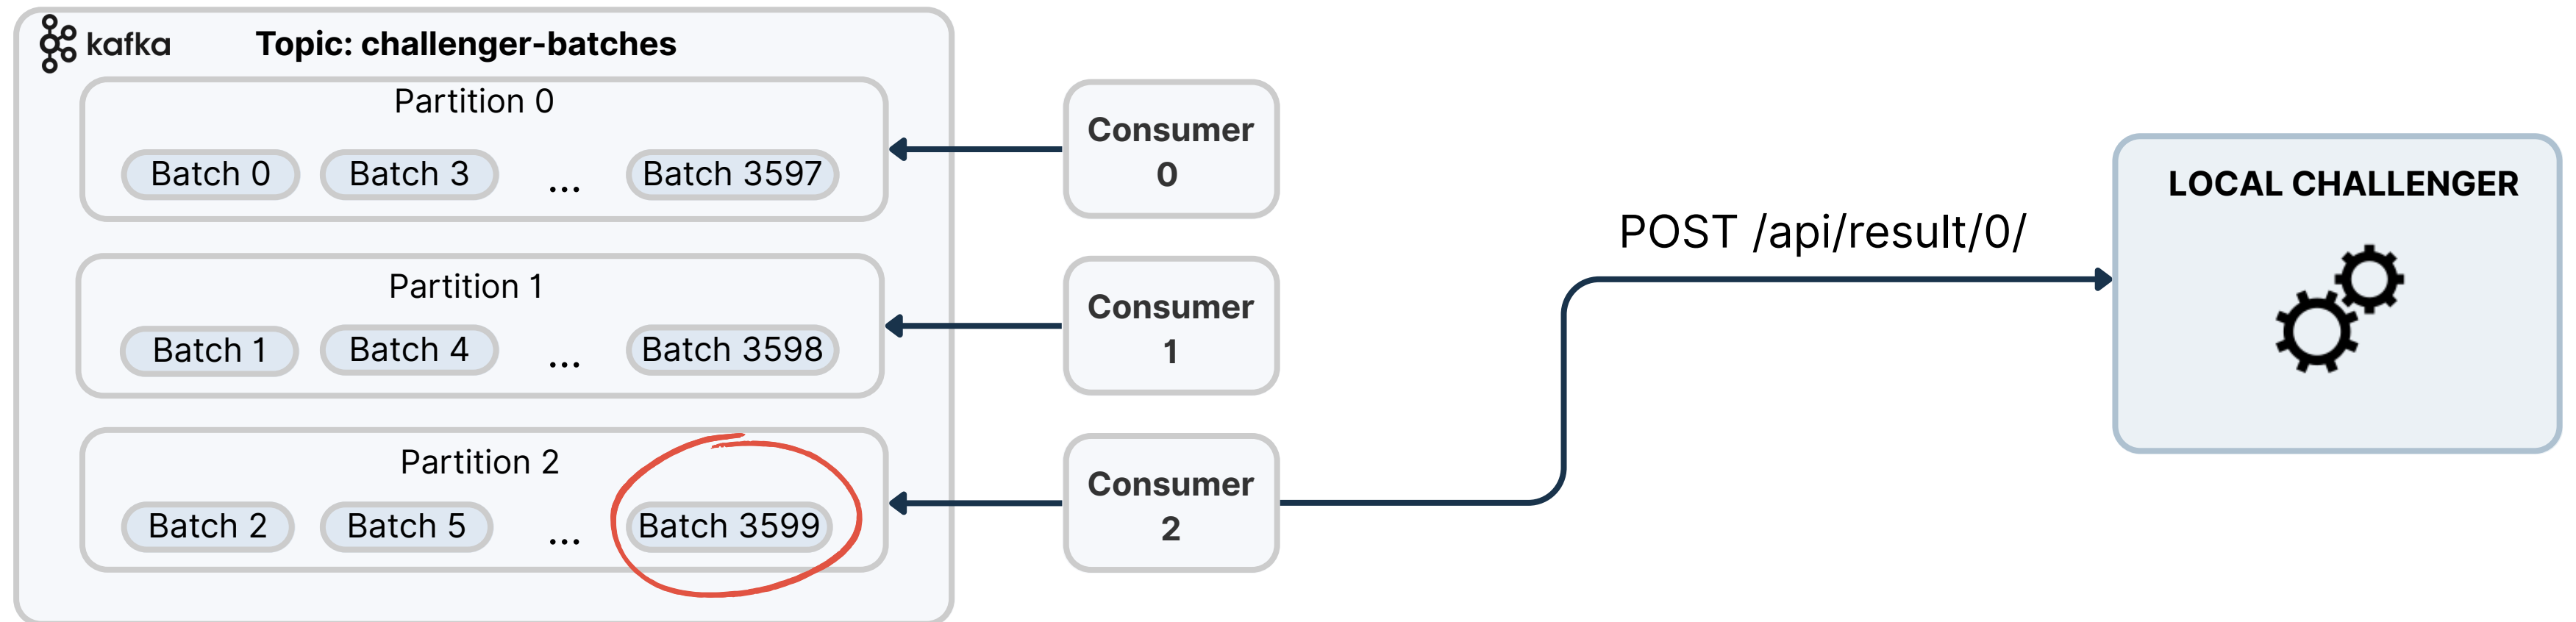
Kafka Producer





Kafka Consumer



Kafka Consumer



Results - Kafka Version




Results - Kafka Version


parallelism ∇	÷	bench_id ∇	÷	throughput ∇	÷	latency_mean ∇	÷	latency_max ∇	÷
1	1	01JZMFGQ81SDJ8G2HEMABQ4T6N		50.22098903979123		39s107ms178µs730ns		1m9s320ms873µs598ns	
2	2	01JZMFN465CZ738SRTNVA0ZX06		85.74736328028007		23s879ms521µs806ns		39s353ms927µs902ns	
3	3	01JZMFS2Q4HEMZTANZTACPDE91		114.48882427842976		18s67ms809µs109ns		28s472ms43µs119ns	
4	4	01JZMFVF1RK6CD4398GS0J6GN7		119.3848778296372		15s743ms100µs850ns		27s804ms834µs136ns	
5	5	01JZMFXW1JKZNNT6MKXYRYNNMJ		149.59246634095877		13s889ms424µs369ns		21s303ms15µs344ns	
6	6	01JZMG4MMMFJFSKHVFHVB6JFT		166.21358244595234		12s775ms733µs853ns		18s912ms902µs556ns	
7	7	01JZMGFHT26P8B5P8FB4WV0S7		178.61263541897438		12s268ms572µs75ns		17s726ms551µs491ns	
8	8	01JZMGQ7AV1BK50KVZVACMQ6TH		175.17665642473474		11s889ms744µs639ns		18s120ms278µs394ns	

Results Query - Kafka Version


Throughput



Max Latency



Avg Latency



So, who has the best performance?



So, who has the best performance?



Conclusions

It should be noted that Flink uses the proprietary Backpressure mechanism, which ensures that the performance of individual queries is the same.

With Kafka, we don't have this mechanism natively and this determines the difference in performance of individual queries.

The difference in performance between the two versions is justified by the intrinsic differences of the two systems.

It should be noted that the collected performance was obtained on a machine with an Intel i9-13980HX and 16 GB RAM in a fully containerised environment, without the use of GPU acceleration, where only the implementation code was executed.



**Thanks for the
Attention!**